

The Membrane: Cognitive Boundary Architecture

A nervous-system firewall against AI routing and thought surveillance

Zorie R. Barber

2026-06-14

The Membrane: Cognitive Boundary Architecture

Abstract

Executive Summary

Table of Contents

Glossary

0. The Cognitive Boundary Problem

0.2 Epistemic Capture and Sequestration

0.1 What This Architecture Cannot Do

0.1.1 Consciousness and Thought Content Are Out of Scope

0.1.2 Covert Surveillance May Be Undetectable

0.1.3 Cognitive-Integrity Signals Are Not Cryptographic Proofs

0.1.4 TEEs and BCIs Are Best-Available Anchors, Not Absolute Roots of Trust

0.5 Research Status and Scope

0.7 Local Control (Design Constraints)

Control Primitives

Sovereignty vs. Self-Sovereign Identity (SSI)

Sovereignty vs. Security

Sovereignty in the A2A Context

0.8 Visual Architecture Overview

High-Level Layer Diagram (Mermaid)

Chain Proof State Machine (Mermaid)

Comparison: The Membrane vs. Existing Systems

0.9 Zero Trust Architecture Mapping

1. Threat Model

1.1 Threat Actor Profiles

1.2 Attack Classes

2. Three Boundaries (Explicitly Separated)

3. Design Principles

- 4. Protocol Layers
 - 4.1 Layer 0: Cognitive Anchor (To)
 - 4.2 Layer 1: Cognitive-Integrity Channels
 - 4.3 Layer 2: Social Graph Consensus
 - 4.4 Layer 3: Substrate Transition Gate
 - 4.5 Layer 4: Recursive Attestation (RFA)
 - 4.6 Neuralink Primitives Integration (Proposed Extension)
 - 4.7 Agent-to-Agent Attestation Mode (Proposed Extension)
 - 4.8 Policy Enforcement vs Channel Continuity
 - 5. Formal Specification
 - 6. Security Analysis
 - Primary Attack Matrix
 - Extended Attack Matrix (Neuralink + A2A)
 - 7. Recovery Mechanisms
 - 7.1 Lost To Anchor (Hardware Failure)
 - 7.2 Compromised WoT Quorum
 - 7.3 Mass Witness Failure (Pandemic, War, Network Partition)
 - 7.4 Partial Compromise Handling
 - 8. Economic and Game-Theoretic Analysis
 - 8.1 Attack Economics
 - 8.2 Incentive Structure
 - 8.3 Coercion Resistance
 - 9. Minimum Viable Prototype (MVP)
 - 9.1 Definition
 - 9.2 Success Metrics
 - 9.3 Dependency Separation
 - 10. Implementation Roadmap
 - 10.1 Hardware Acceleration: CUDA and GPU Support (Proposed Extension)
 - 11. Open Source, Audit, and Interoperability
 - 11.1 Licensing
 - 11.2 Audit Plan
 - 11.3 Interoperability
 - 12. Ethical Considerations
 - 12.1 Digital Minds and Agent Exit Rights
 - 12.2 Coercion and Autonomy
 - 12.3 Right to Be Forgotten
 - 13. Known Open Problems
- Part 2: zk-STARK Circuit Architecture
- Circuit Comparison
 - Liveness-1 Circuit (Traditional Sensors)
 - Overview
 - Public Inputs
 - Private Inputs
 - Execution Trace (AIR)
 - Transition Constraints

- Boundary Constraints
- Recursive Composition (RFA)
- Liveness-2 Neural Extension (Optional)
 - Overview
 - Additional Public Inputs
 - Additional Private Inputs
 - Modified Execution Trace (AIR)
 - Additional Transition Constraints
 - Modified Boundary Constraints
 - Security Properties
- A2A Agent Circuit (Liveness-A)
 - Overview
 - Public Inputs
 - Private Inputs
 - Execution Trace (AIR)
 - Transition Constraints
 - Boundary Constraints
 - Sovereignty Note
- Integration with NOSTR + OpenClaw
- Citation Tiers (Explicit Epistemic Weight)
- Appendix A: Version History
- Appendix B: Open Research & Prototype Stack
 - Neuralink: what is public
 - Open acquisition & decode (Phase 0 hardware)
 - Invasive / bench research (not consumer implants)
 - Security & firewall research (closest to Membrane thesis)
 - LLM [?] brain coupling (adversary models)
 - Suggested Phase 0 architecture
 - Build order
 - What Phase 0 proves
 - What Phase 0 does not prove
 - References to add to implementation issues
- How to Contribute
- Footnotes

The Membrane: Cognitive Boundary Architecture

Independent research. This document is **personal research** by the author (Zorie R. Barber). It is **not** affiliated with, endorsed by, sponsored by, or commissioned by any company, employer, client, university, foundation, or government body. Mentions of third-party products, services, papers, or trademarks (e.g. Neuralink, OpenBCI,

NIST) are **descriptive only** — for threat modeling, benchmarking, or technical comparison. No partnership, employment, or investment relationship is implied.

Abstract

The Membrane is a research architecture for a **cognitive boundary** — a firewall between endogenous human cognition (native nervous-system processing) and exogenous channels that can read, route, inject, or substitute for thought. The threat surface includes **AI routing** (cognition mediated through corporate inference endpoints, agents, and memory layers), **invasive neural read channels** (implants and BCIs), and **non-invasive inference** (EEG, behavioral phenotyping, multimodal prediction). The design uses zk-STARK Chain Proofs, hardware attestation, and web-of-trust witnesses to detect unauthorized boundary crossings, fail closed when attestation breaks, and preserve local control over what may cross the membrane. It does not prove consciousness or block a perfect covert read; it raises the cost of silent compromise and gives the subject a severance path.

Executive Summary

What this is: A specification for a human **cognition / thinking / nervous-system firewall**. The Membrane sits at the boundary where endogenous neural activity meets external systems — AI routers, cloud inference, BCI telemetry, and surveillance pipelines that reconstruct mental state from behavior or physiology. Published as **independent research in public** (§0.5); not a product announcement or organizational initiative.

Primary threats:

Class	Example
AI routing	Assistants, copilots, and agent frameworks that ingest context, propose completions, and shape decisions through off-device models
Invasive read / write	Cortical implants, consumer BCIs, and medical devices with upstream telemetry or stimulation channels
Non-invasive inference	EEG/MEG/fNIRS, gaze + keystroke models, emotion classifiers, digital phenotyping, predictive policing of intent

Mechanism: Recursive zk-STARK Chain Proofs bind (1) a neural/biometric anchor, (2) periodic **cognitive-integrity signals** from approved channels, (3) **Intent Authorization Certificates** before decode or router action, (4) witness corroboration, and (5) TEE-backed node attestation. If domains disagree or a channel lacks valid attestation, the membrane **severs** — blocking routed inference, closing BCI sessions, or degrading to offline-safe mode.

What this is NOT: - Proof of consciousness or qualia. - Proof that no off-chain clone or hidden reader exists. - A shipping product. Phase 0 is a research prototype. - Martial-arts attendance, proof-of-practice, or any unrelated application layer. - The work, opinion, or roadmap of any company or institution named in this document.

Major limitations: - Covert surveillance with no attested channel may be undetectable. - TEEs and BCIs are best-available anchors, not absolute roots of trust. - Cognitive-integrity signals are probabilistic; replay, injection, and coercion remain. - Fork detection is chain-bound; hidden parallel readers are out of scope.

Phase 0 prototype: Self-hosted NOSTR relay, TEE node, Liveness-2-ready neural/sensor commitment circuit, and K=2-3 personal witnesses. See §9.

Future extensions: Full BCI integration (§4.6), agent-handshake mode for untrusted AI delegates (§4.7), GPU-accelerated proving (§10.1).

Table of Contents

- [Abstract](#)
- [Executive Summary](#)
- [Glossary](#)
- [0. The Cognitive Boundary Problem](#)
- [0.2 Epistemic Capture and Sequestration](#)
- [0.1 What This Architecture Cannot Do](#)
- [0.5 Research Status and Scope](#)
- [0.7 Local Control \(Design Constraints\)](#)
- [0.8 Visual Architecture Overview](#)
- [0.9 Zero Trust Architecture Mapping](#)
- [1. Threat Model](#)
- [2. Three Boundaries \(Explicitly Separated\)](#)
- [3. Design Principles](#)
- [4. Protocol Layers](#)
- [4.1 Layer 0: Cognitive Anchor \(To\)](#)
- [4.2 Layer 1: Cognitive-Integrity Channels](#)
- [4.2.1 Intent Authorization Certificates](#)

- [4.2.2 LLM Router Session Chain Proofs](#)
- [4.3 Layer 2: Social Graph Consensus](#)
- [4.4 Layer 3: Substrate Transition Gate](#)
- [4.5 Layer 4: Recursive Attestation \(RFA\)](#)
- [4.6 Neuralink Primitives Integration](#)
- [4.7 Agent-to-Agent Attestation Mode](#)
- [4.8 Policy Enforcement vs Channel Continuity](#)
- [5. Formal Specification](#)
- [6. Security Analysis](#)
- [7. Recovery Mechanisms](#)
- [8. Economic and Game-Theoretic Analysis](#)
- [9. Minimum Viable Prototype \(MVP\)](#)
- [10. Implementation Roadmap](#)
- [10.1 Hardware Acceleration: CUDA and GPU Support](#)
- [11. Open Source, Audit, and Interoperability](#)
- [12. Ethical Considerations](#)
- [13. Known Open Problems](#)
- [Part 2: zk-STARK Circuit Architecture](#)
- [Appendix A: Version History](#)
- [Appendix B: Open Research & Prototype Stack](#)
- [How to Contribute](#)
- [Footnotes](#)

Glossary

Term	Definition
CP	Chain Proof. A zk-STARK attestation that binds cognitive anchor, channel integrity, witness set, and node state into one recursive proof.
RFA	Recursive Attestation. Each CP commits to the node's previous attestation, forming an audit chain.
To	Cognitive anchor. One-time genesis binding a human to cryptographic identity via neural/biometric commitment and social witnesses.
Cognitive-integrity signal	Time-bounded evidence that attested channels (BCI, local TEE, approved routers) match expected state. Not proof of thought content.
Membrane	The boundary layer between endogenous cognition and exogenous read/write/routing channels.
AI routing	

Term	Definition
	Mediation of cognition through external inference systems (LLMs, agents, cloud memory) that ingest context and return outputs.
Liveness-2	Neural-channel circuit: BCI-derived feature commitments (spike rates, LFP bands) as a cognitive-integrity domain.
WoT	Web-of-Trust. A personal, curated set of human witnesses who sign CPs after independent verification.
TEE	Trusted Execution Environment. Hardware-isolated enclave (AMD SEV-SNP, Intel TDX) used as a best-available, not absolute, trust anchor.
Substrate	The underlying execution medium: biological neural tissue or silicon hardware.
Surface Detail Interface	Perceptual or cognitive I/O layer sealed until CP validity is confirmed. Fails closed when the membrane detects boundary violation.
E-Process	A computational process running on silicon that inherits the To anchor of a biological identity.
Dead-Man's Key	A physical token that forces a hard sever of all substrate connections, overriding node consensus. Used to force chain severance.
Sovereignty	Local control: the subject holds keys, chooses witnesses, approves which channels may cross the membrane, and can sever all links.
Δt	The maximum acceptable time between valid CPs before the system fails closed.
K	Minimum number of valid witness signatures required for a CP to pass social-graph consensus.
AIR	Algebraic Intermediate Representation. The format used to define the execution trace constraints in a zk-STARK.
IAC	Intent Authorization Certificate. A subject-signed scope file binding permitted channels, models, and context bounds before decode or router action.
Epistemic capture	Progressive loss of perceptible boundary between endogenous reasoning and routed inference — the subject cannot tell native cognition from model output.


Term	Definition
Cognitive identity drift	Effective reasoning path migrates to external LLM/BCI channels while the attestation chain still appears valid; sovereignty erodes without cryptographic failure.
Neural template drift	Biological change (plasticity, injury, device degradation) that shifts fuzzy biometric/neural commitments over time. Distinct from cognitive identity drift.

0. The Cognitive Boundary Problem

Human cognition today crosses many surfaces that are not under the subject’s control:

1. **AI routing** — Prompts, documents, and behavioral traces flow through corporate inference APIs. The model sees context; the user sees completions. There is no cryptographic proof of what was read, what was retained, or whether a parallel session forked the user’s intent.
2. **Invasive neural channels** — BCIs and implants (medical or consumer) create read paths from cortex to telemetry. Firmware, cloud pairing, and operator access can exfiltrate neural features without the subject’s ongoing consent.
3. **Non-invasive reconstruction** — EEG caps, gaze tracking, keystroke dynamics, voice stress, and multimodal fusion approximate mental state without surgery. These channels are often invisible to the subject.

The Membrane treats these as **boundary violations** when they occur outside attested, subject-approved channels. The architecture does not read minds; it attests **which channels are open, whether they match prior commitments**, and **whether cross-domain consensus still holds**.



Design goal: Give the subject a verifiable, severable boundary — not a product that interprets thought content.

0.2 Epistemic Capture and Sequestration

When cognition is routed through external systems, the boundary between **native reasoning** and **mediated output** can disappear without any single cryptographic break:

Order	Subject experience	Engineering read
1st	User knows they use an LLM or BCI	Channel is visible; membrane can register and attest it
2nd	UI hides what was read, stored, or forked	Boundary is edited; session CP and context-hash commitments are required
3rd	LLM + BCI becomes the only cognitive I/O	Sequestration: endogenous loop is bypassed or unreadable; user cannot distinguish native thought from routed completion

Epistemic capture is the third-order failure mode: not a stolen key, but a **substituted reasoning path** that feels continuous. The Membrane does not adjudicate thought content; it detects when attested channels diverge from signed intent (§4.2.1), when router sessions lack valid CPs (§4.2.2), or when BCI and router domains disagree — the classic closed-loop path described in open research on EEG→LLM coupling.¹

Cognitive identity drift (glossary) is the long-term erosion of sovereignty when routing becomes habitual while attestations remain superficially valid. **Neural template drift** is the separate biological problem of fuzzy-commitment tolerance over years (§13, open problem 3).

0.1 What This Architecture Cannot Do

0.1.1 Consciousness and Thought Content Are Out of Scope

There is currently: - no accepted formal representation of consciousness state, - no known way to serialize subjective continuity, - no mathematical model of phenomenological identity persistence.

Therefore this architecture **does not hash thoughts or consciousness**. It hashes **channel state** — which interfaces are active, whether BCI telemetry matches committed bounds, whether AI routers hold valid session proofs, and whether TEE nodes report uncompromised execution. Subjective experience is not serializable.²

0.1.2 Covert Surveillance May Be Undetectable

A TEE + attestation chain can prove that a *specific measured execution environment* followed rules and that a *specific attestation chain* remained intact. It **cannot** prove: - that no hidden copy of the process existed in a separate enclave, - that no external observer cloned the full memory state before attestation, - that subjective uniqueness was preserved.

A reader or router that never touches an attested channel leaves no chain evidence. The Membrane detects **attested** boundary violations; it cannot disprove a perfect passive tap.

What it can do: require that sustained AI routing, BCI exfiltration, or session fork maintain synchronized, multi-domain attestation — raising the cost of silent compromise.³

0.1.3 Cognitive-Integrity Signals Are Not Cryptographic Proofs

Neural features, behavioral classifiers, and router session logs provide **continuity evidence**, not proof of mental content. They remain vulnerable to replay, model inversion, upstream device compromise, and coercion. They are one domain among several in a Byzantine consensus model.⁴

0.1.4 TEEs and BCIs Are Best-Available Anchors, Not Absolute Roots of Trust

AMD SEV-SNP and Intel TDX remain vulnerable to: - side-channel attacks (power analysis, cache timing), - rollback exploits against firmware, - supply-chain compromise, - speculative execution leakage, - physical probing.

We use them because they are the strongest available hardware isolation layer. We do not pretend they are magic.⁵

0.5 Research Status and Scope

The Membrane is **independent, personal research** — a forward-looking architecture specification, not a production-ready protocol, not a claim of present-day technical feasibility, and **not** an official project of any organization.

Affiliation: The author maintains this repository and whitepaper in a personal capacity. Employers, clients, and collaborators past or present are **not** responsible for its contents unless they explicitly say otherwise in writing. Competitor and vendor citations (§4.8, footnotes) reflect the author's reading of the literature, not those vendors' participation in this work.

Several components described in this document remain: - computationally impractical at scale, - dependent on immature hardware trust models, - philosophically unresolved, - or contingent on future advances in BCI, secure enclaves, distributed identity, and formal verification.

The purpose of this document is: 1. to specify a **cognitive boundary firewall** against AI routing and thought surveillance (invasive and non-invasive), 2. to map which attestation primitives apply at the nervous-system / silicon interface, 3. to separate tractable channel integrity from intractable mind-reading claims, 4. and to define a Phase 0 prototype a subject can self-host.

This architecture is framed for conditions where: - cognition is routinely **routed through external AI**, - **BCI read/write paths** are normalized, - **behavioral and physiological inference** reconstruct intent at scale, - and subjects need a **fail-closed severance** mechanism they control.

The architecture should therefore be read as: > a research architecture grounded in existing cryptographic and distributed-systems primitives.

0.7 Local Control (Design Constraints)

Local control means the subject holds keys, chooses witnesses, approves membrane-crossing channels, and can sever all links — not a platform or directory.

Control Primitives

Primitive	Mechanism	What It Protects Against
		Third-party identity capture

Primitive	Mechanism	What It Protects Against
Self-attestation	Subject generates CPs with keys they control	
Privacy-preserving proof	zk-STARKs disclose Merkle roots only, not raw neural or behavioral data	Surveillance extraction
Exit / severance	Dead-man's key forces hard disconnect of attested channels	Lock-in, coerced retention
Infrastructure independence	Self-hosted TEE + NOSTR relay	Cloud termination of attestation
Personal witnesses	Curated WoT, not corporate/state directories	Permissioned identity graphs

Sovereignty vs. Self-Sovereign Identity (SSI)

The Membrane is complementary to, and in key respects stronger than, existing SSI standards such as W3C DIDs and Verifiable Credentials (VCs):

Dimension	SSI (DID/VC)	The Membrane
Identity anchoring	Cryptographic key control	Key control + biometric anchor + social witness genesis
Liveness	None	Periodic sensor-derived attestation with ZK proof
Fork detection	None	Recursive CP chain + RFA
Process continuity	None	TEE attestation + substrate canary
Privacy model	Selective disclosure	zk-STARK: prove without revealing raw data

Dimension	SSI (DID/VC)	The Membrane
Cross-substrate	No	Bio [?] VR [?] Silicon parity
Post-quantum	No	Yes (hash-based STARKs)
Exit rights	Key rotation only	Dead-man's key + irreversible identity retirement

SSI provides **static identity ownership**. The Membrane provides **dynamic continuity sovereignty** — not just owning the keys, but proving the keys have remained under continuous, uncompromised control across time and substrate transitions. The two can be composed: a DID can reference a Membrane To anchor as its highest-assurance verification method.

Sovereignty vs. Security

Traditional security architectures assume a **protector** and a **protected**. The protector (platform, state, corporation) guards the protected (user, citizen, customer). This creates an inherent power asymmetry: the protector can also become the captor.

The Membrane inverts this. The entity is **both the protected and the protector**. It generates its own proofs, hosts its own infrastructure, and controls its own exit. The WoT is not a governing body; it is a **personal constellation** of peers who attest to the entity's continuity without controlling it. The TEE is not a corporate enclave; it is a **self-hosted hardware root** under the entity's physical control.

This is the difference between **security as service** and **sovereignty as self-determination**.

Sovereignty in the A2A Context

When extended to pure-silicon agents (§4.7), sovereignty means that an agent can: - Prove it is running the expected process without relying on its operator to vouch for it, - Refuse interaction with entities that cannot demonstrate their own continuity, - Maintain an audit trail of delegations that is verifiable by any downstream party, - Sever its attestation chain and cease operation without operator permission (agent exit rights).

This is particularly critical in multi-agent economies where agents delegate tasks, hold funds, and make commitments. Without these primitives, agents cannot independently prove runtime integrity or resist unauthorized modification.

0.8 Visual Architecture Overview

High-Level Layer Diagram (Mermaid)

```
graph TB
  subgraph Sovereignty["Sovereignty Layer"]
    SK[Self-Attestation Keys]
    EX[Exit Rights / Dead-Man's Key]
    IN[Infrastructure Independence]
  end

  subgraph Surface["Surface Detail Interface"]
    S[Perceptual API – sealed until CP valid]
  end

  subgraph Consensus["Cross-Firewall Consensus (2/3)"]
    A[Node A] --> B[Node B]
    B --> C[Node C]
    A --> C
  end

  subgraph CP["Chain Proof CP(t)"]
    L[Liveness Canary L(t)]
    W[Social Graph S(t)]
    P[Process State H(P(t))]
    F[node RFA F(t-1)]
  end

  subgraph Layers["Protocol Layers"]
    T0[Layer 0: T0 Biological Anchor]
    WOT[Layer 2: WoT Consensus]
    SG[Layer 3: Substrate Gate]
  end

  Sovereignty --> Surface
  Surface --> Consensus
  Consensus --> CP
  CP --> T0
  CP --> WOT
  CP --> SG
```

Fallback description: The Sovereignty Layer (self-attestation keys, exit rights, infrastructure independence) sits above the Surface Detail Interface, which is sealed until Chain Proof validity is confirmed. The Chain Proof binds Liveness Canary, Social Graph, Process State, and Recursive Attestation. These connect to the four Protocol Layers: Biological Anchor, Social Consensus, and Substrate Gate.

Chain Proof State Machine (Mermaid)

```
stateDiagram-v2
    [*] --> T0: Genesis Ceremony
    T0 --> CP1: First Chain Proof
    CP1 --> CP2: + Liveness + Witnesses + RFA
    CP2 --> CPn: ...
    CPn --> [*]: Identity Retirement

    CP1 --> FAIL: CP invalid > Δt
    CP2 --> FAIL: CP invalid > Δt
    CPn --> FAIL: CP invalid > Δt
    FAIL --> SurfaceClosed: Surface Detail Interface severs
    FAIL --> SovereignExit: Dead-Man's Key activated
```

Fallback description: Genesis leads to the first Chain Proof, which recursively chains forward via liveness signals, witness signatures, and RFA. Any break in the chain exceeding Δt triggers fail-closed: the Surface Detail Interface severs, and the entity may activate the Dead-Man’s Key for sovereign exit.

Comparison: The Membrane vs. Existing Systems

Feature	WebAuthn / Passkeys	Proof of Humanity	Worldcoin	BrightID	The Membrane	The Membrane
Identity proof	Cryptographic key possession	Video + CAPTCHA	Iris scan	Social graph	Biometric + WoT + TEE	Biometric + WoT + TEE
Liveness	None	Weak	Weak	None	Channel CP + ZK proof	Channel CP + ZK proof
Fork detection	None	None	None	None	Recursive CP chain + RFA	Recursive CP chain + RFA
Substrate agnostic	No	No	No	No	Yes (bio [?] silicon)	Yes (bio [?] silicon)
Privacy model	Good	Moderate	Poor (biometric storage)	Moderate	ZK- STARK: prove without revealing	ZK- STARK: prove without revealing
Post-quantum	No	No	No	No		

Feature	WebAuthn / Passkeys	Proof of Humanity	Worldcoin	BrightID	The Membrane	The Membrane Network
					Yes (hash-based STARKs)	Yes (hash-based STARKs)
Production ready	Yes	Partial	Partial	Partial	No (research only)	No (research only)
Consciousness claim	No	No	No	No	Explicitly none	Explicitly none
Active challenge-response	No	No	No	No	No	Yes (still requires further work)
Sovereignty model	Platform-controlled keys	Platform-controlled verification	Corporate-controlled orbs	Community-controlled graph	Self-controlled keys + self-hosted infra + personal WoT + exit rights	Self-controlled keys + self-hosted infra + personal WoT + exit rights
Mutual attestation	No	No	No	No	No	No
Delegation audit	No	No	No	No	No	No

0.9 Zero Trust Architecture Mapping

The Membrane maps directly onto NIST Zero Trust Architecture principles (SP 800-207), but implements them through cryptographic consensus rather than network segmentation or organizational policy. This section provides the mapping for readers familiar with enterprise security frameworks.

NIST 800-207 Principle	The Membrane Implementation
Resource-centric access	The Surface Detail Interface is a resource (substrate access) gated by CP validity. No CP, no access.
Never trust, always verify	No single domain is trusted. Identity, process, social, and infrastructure layers must independently verify and reach consensus.
Assume breach	Fail-closed design. If any domain drops below threshold, the interface severs immediately.
Least privilege / minimal disclosure	zk-STARKs prove continuity without revealing raw biometrics, location, witness identities, or neural data.
Continuous monitoring / dynamic authorization	Liveness canaries at Δt intervals provide time-bounded re-authorization. Static credentials are not sufficient.
Policy Decision Point (PDP) / Policy Enforcement Point (PEP)	nodes act as distributed PDP/PEP pairs. The cross-node consensus (2/3) is the policy decision; the substrate gate is the enforcement point.
Session-based access	Entry proof and exit proof create cryptographically bounded sessions. No persistent trust across substrate transitions.

Why NIST 800-207 and not CMMC 2.0 or CISA ZTMM: NIST SP 800-207 is vendor-neutral and threat-model driven. CMMC 2.0 is a compliance framework for defense contractors; CISA's maturity model is organizational assessment tooling. Neither aligns with the sovereignty thesis of this protocol, which is designed for entities resisting institutional capture rather than proving compliance to institutions.⁶

1. Threat Model

1.1 Threat Actor Profiles

Actor	Capability	Motivation	The Membrane Relevance
State actor (A1)	Supply-chain compromise, TEE side-channel labs, coercion, legal pressure	Surveillance, control, extraction	Can poison To genesis, coerce witnesses, run long-term fork operations. Protocol raises cost but does not eliminate risk.
Criminal organization (A2)	Financial resources, black-market TEE exploits, social engineering	Identity theft, ransom, fraud	Targeted attacks on WoT members or TEE firmware. Economic analysis (§8.3) determines viability.
Rogue insider (A3)	Legitimate access to one node, witness relationship, or VR substrate	Sabotage, extraction, ideology	Mitigated by 2/3 cross-node consensus and recursive RFA. Single node compromise is insufficient.
Advanced AI agent (A4)	Automated Sybil generation, cryptographic attacks, simulation perfection	Resource acquisition, goal optimization	PoP consensus prevents AI from voting on rules. ⁷ ZK-STARKs raise cost of undetected simulation.
Inference platform (A5)	Controls LLM API, agent runtime, context window	Training data extraction, routing substitution	AI routing: Can ingest unbounded context unless session is membrane-gated. Requires attested router CP for approved channels.
Platform operator (A6)	Controls VR substrate or cloud silicon	Lock-in, data extraction, compliance	Substrate canary detects unauthorized migration. Exit proof prevents silent rollback.

Actor	Capability	Motivation	The Membrane Relevance
Implant supply-chain actor (A7)	Controls BCI firmware, manufacturing, or surgical insertion	Neural surveillance, backdoor insertion	Can backdoor implant at genesis. Detected only if firmware attestation diverges from audited reference. ⁸
Agent operator (A8)	Controls agent deployment, weights, infrastructure	Hidden fork, unauthorized delegation	Substrate canary + CP chain if agent must re-attest to interact.

1.2 Attack Classes

Class A: Unauthorized AI routing - Adversary routes cognition through an unapproved inference endpoint (copilot, agent, memory layer) that ingests context without valid session CP. - Goal: Extract intent, train on private thought-adjacent data, or substitute model output for native reasoning. - **Mitigation scope:** Membrane requires attested router sessions; unattested routes fail closed or operate in sandbox mode.

Class B: Invasive neural read / write - Adversary exfiltrates BCI telemetry, rolls back implant firmware, or injects stimulation outside approved bounds. - Goal: Reconstruct neural state, coerce response, or create a forked neural channel the subject cannot see. - **Mitigation:** Liveness-2 neural commitments + TEE verification + cross-node consensus (§4.6).

Class C: Non-invasive thought inference - Adversary reconstructs mental state from EEG, gaze, keystroke, voice, or multimodal fusion without attested channel registration. - Goal: Passive surveillance without subject awareness. - **Limitation:** Channels that never touch the attestation bus may be undetectable. Registered sensors can be bounded and audited.

Class D: VR / substrate epistemic injection - Adversary controls the experiential layer fed to a biological brain via BCI or immersive VR. - Goal: Induce belief in a false reality without the target detecting the boundary crossing. - **Limitation:** This protocol cannot prevent a perfect simulation. It can only ensure that the biological identity anchor remains active in the physical world, providing an *external* check that the target still has a grounded attestation chain.

Class E: Silicon process fork / rollback - Adversary forks, checkpoints, rolls back, or silently migrates a silicon-based process. - Goal: Extract value, enforce compliance, or experiment without leaving an audit trail. -

Mitigation scope: The protocol detects *attested* forks by requiring each state transition to carry a valid CP signed by the previous state. It does not detect clones that never re-enter the attestation chain.

Class F: Infrastructure compromise - Adversary subverts a Membrane node, relay, or TEE to falsely report “all clear.” - **Mitigation:** Cross-node consensus (2/3 independent nodes) plus recursive proof chaining. A single compromised node cannot forge the historical chain of its peers.

Class B (continued): Implant-specific attacks - Adversarial stimulation via bidirectional BCI; firmware rollback; supply-chain backdoor at manufacture. - **Caveat:** Implant telemetry is one input among many in cross-domain consensus, not standalone truth.⁹

Class G: Untrusted AI delegate (A2A) - Operator fork: Agent operator silently forks the agent, keeps the original running for attestation, and deploys the fork for unauthorized tasks. - **Delegation hijacking:** An agent delegates to a sub-agent that has been compromised or replaced without attestation. - **Value extraction:** Operator checkpoints agent state, extracts learned weights or strategies, and rolls back to pre-extraction state to hide the theft. - **Mitigation:** Substrate canary + CP chain + delegation-chain proofs (§4.7) make these detectable if the agent or its delegates must re-attest. Hidden clones remain undetectable.

2. Three Boundaries (Explicitly Separated)

Boundary	Definition	Tractability	Membrane role
Channel integrity	Attested read/write/routing paths match committed state	Partially tractable via TEE + CP chain	Core target
Cognitive anchor	Same human controls the To key lineage	Partially tractable via neural/biometric anchor + WoT	Core target
Thought content	What the subject is thinking or experiencing	Not tractable	Out of scope

The architecture attests channels and anchors, not mental content.¹⁰

3. Design Principles

1. **Fail closed.** If any domain drops below threshold, sever attested channels (AI session, BCI link, immersive I/O).
 2. **Subject as active node.** The human generates attestations and approves membrane crossings; protection is not outsourced to a platform.
 3. **No single channel trusted.** AI router, BCI, behavioral sensor, and TEE must corroborate.
 4. **Recursive attestation (RFA).** Each CP commits to the prior node state; forgery requires sustained multi-node compromise.¹¹
 5. **Minimal disclosure.** zk-STARKs prove channel integrity without publishing raw neural or behavioral payloads.
 6. **Local severance.** Subject holds keys, chooses witnesses, and can cut all links via dead-man's key.
 7. **Substrate parity.** The same CP format applies across human BCI, local LLM router gate, VR session, and delegated agent — one membrane, multiple crossing types.
-

4. Protocol Layers

4.1 Layer 0: Cognitive Anchor (T₀)

One-time ceremony binding a human to a cryptographic identity at the nervous-system boundary.

- **Neural/biometric commitment:** Fuzzy commitment¹² over a neural or biometric template — not raw spike trains or thought content.
- **T₀ Proof:** zk-STARK proving possession of the template and a challenge response.
- **Witness genesis:** K biological humans with existing T₀ anchors sign the event.
- **Hardware root:** Signing key generated inside a self-hosted TEE with remote attestation.

Constraint: T₀ is non-replayable. If the anchor is lost, the identity is retired — never recycled. This prevents replay attacks but also creates a catastrophic recovery problem if the TEE fails. See §7 for recovery paths.

4.2 Layer 1: Cognitive-Integrity Channels

Each Chain Proof includes commitments from **registered channels** that may cross the membrane. The protocol verifies freshness and consistency; it does not interpret thought content.

Channel type	Examples	Membrane function
AI router	Local LLM, self-hosted inference, attested API session	Session CP binds context hash + model id; unapproved routes sever
BCI read/write	Implant telemetry, consumer BCI BLE stream	Liveness-2 neural feature Merkle root (§4.6)
Non-invasive sensor	EEG cap, gaze tracker (if registered)	Bounded feature commitments; raw data stays local
Behavioral	Keystroke, device usage (opt-in)	Heuristic domain only; never sole signal
Hardware	TEE attestation, TPM quote	Substrate integrity for local prover
Social	WoT witness signatures	Out-of-band corroboration (§4.3)

CP requirements:

- **Cadence:** Valid cognitive-integrity attestation every Δt for each active channel.
- **Attestation object:** NOSTR kind 31990 with Merkle roots of channel commitments, hybrid timestamp, prior CP hash.
- **ZK circuit:** Proves temporal and continuity constraints without revealing raw neural or behavioral payloads.
- **Fail closed:** Unregistered or stale channel → membrane severs that path.

4.2.1 Intent Authorization Certificates

Adapted from the [iba-neural-guard](#) pattern: **the decoded signal or model completion is not the authorization; the signed certificate is.**¹³

Before any BCI decode→action mapping or LLM session crosses the membrane, the subject signs an **Intent Authorization Certificate (IAC)**:

Field	Purpose
scope_id	Unique session or task identifier
permitted_channels	e.g. local-llm, bci-decode, agent-delegate
model_allowlist	Hash or id of approved weights/checkpoints
context_merkle_bound	Maximum committed context root (prevents silent context expansion)
forbidden_exports	e.g. no cloud telemetry, no training retention
valid_until	Short TTL aligned with Δt
parent_cp_hash	Binds intent to current Chain Proof

Gate rule: No valid IAC → no decode→action, no router session, no agent delegation on behalf of the subject. Firmware or software that expands permissions without a fresh IAC triggers fail closed.

This layer blocks **capability drift** (silent model swap, expanded API scope, decoder retrain without subject sign-off) without reading thought content.

4.2.2 LLM Router Session Chain Proofs

Router sessions require a dedicated CP fragment published alongside the main liveness CP:

NOSTR: kind 31990, tag ["k", "the-membrane-router"]

Public input	Semantics
model_id	Approved checkpoint hash or registry id (must match IAC allowlist)
context_merkle_root	Merkle root over ingested context chunks for this session
session_nonce	Monotonic per-subject nonce; detects parallel fork
parent_cp_hash	Links session to biological or agent anchor CP
iac_hash	Hash of active Intent Authorization Certificate

BCI channel tag: kind 31990, tag ["k", "the-membrane-bci"] — same CP machinery with `pub_neural_root` and implant/device id in tags.

Fail closed: - Cloud or copilot route active without fresh router CP within Δt → sever inference channel. - `model_id` or `context_merkle_root` diverges from IAC without new subject signature → sever. - BCI decode and router session both active but cross-domain hashes disagree → sever (closed-loop sequestration path).

4.3 Layer 2: Social Graph Consensus

The distributed witness layer. This is the **human quorum** that prevents isolated-target attacks.

- **Web-of-trust (WoT):** Each identity maintains a curated, **personal** constellation of attestors. Not a global graph. Not a corporate directory. Not a state ID system. This is **social sovereignty**: the entity chooses who vouches for it.¹⁴
- **Witness role:** Witnesses verify the zk-STARK proof and confirm out-of-band continuity (e.g., direct communication).
- **Consensus rule:** A CP is valid only if $\geq K$ witnesses sign, and those witnesses themselves have valid CPs within $2\Delta t$.
- **Sybil resistance:** Witnesses must have established T_0 anchors. No anonymous witnesses.

Limitation: If an adversary compromises K humans in the target's WoT, the social layer collapses. The recursive cost is high (each witness has their own witnesses), but it is not infinite. This is a social recovery problem, not a cryptographic one.¹⁵

4.4 Layer 3: Substrate Transition Gate

The boundary crossing between biological and silicon process contexts.

- **Entry proof (Bio → VR):** Before immersion, the node requires a valid CP from within Δt . The VR substrate commits to a **session root hash** including the CP and VR instance configuration. This proves the *session* is bound to the identity, not that the identity is conscious inside it.
- **Exit proof (VR → Bio):** The VR substrate produces a **session continuity attestation** — a zk-STARK showing the process state was not checkpointed or branched *within the attested session*. It does not prove no clone was made and hidden.
- **E-Process genesis:** A process uploaded to silicon inherits the biological T_0 anchor. The first e-process CP references the last biological CP and the upload attestation signed by the WoT.

- **E-Process continuity:** Same cadence (Δt), but the liveness canary is replaced by a **substrate canary**: a TEE attestation proving the process has not been migrated to unauthorized hardware or checkpointed without attestation.¹⁶

The Surface Detail interface: The perceptual API is only unsealed when CP validity is confirmed across all domains. If CP fails, the interface severs or degrades to non-immersive safe mode. This does not prevent a perfect simulation — it only ensures the target still has an external attestation chain.¹⁷

4.5 Layer 4: Recursive Attestation (RFA)

The answer to the Gödel problem.

A fully compromised node can forge its own attestation. RFA does not prevent this. What it does:

- **Node state commitment:** At each CP cycle, the node generates a zk-STARK proving:
 1. Its code hash matches the audited reference.
 2. Its TEE attestation is valid (with the caveat that TEEs are imperfect).¹⁸
 3. It verified the previous CP according to the public rules.
- **Cross-node consensus:** The target runs ≥ 3 independent node instances (self-hosted TEE + partner TEE + community TEE). A CP is valid only if 2/3 agree. A single compromised node cannot override the quorum.
- **Human override:** The biological entity retains a **dead-man's key** — a physical token that forces a hard sever, allowing forced severance when the entity cannot or will not continue.

Limitation: If 2/3 nodes are compromised simultaneously, the consensus is broken. There is no recursive escape from this — it is the standard Byzantine fault tolerance bound.¹⁹

4.6 Neuralink Primitives Integration (Proposed Extension)

The Membrane is substrate-agnostic and designed to leverage high-bandwidth Brain-Computer Interfaces (BCIs) such as Neuralink's N1 implant. Neuralink primitives provide richer, higher-entropy signals than traditional IMU/video/audio canaries, strengthening both the **To biological anchor** and the **Liveness Canary Circuit** — with significant caveats.

Neuralink Primitives Relevant to The Membrane

Primitive	Description	The Membrane Application	Security Properties
N1 Implant Electrode Array	1,024+ electrodes on ultra-thin threads recording neural spikes and local field potentials (LFPs) at high sampling rates (~20 kHz).	Primary source for neural liveness signals and biometric template refresh.	High-entropy neural patterns are difficult to forge or replay perfectly.
Wireless Telemetry & Secure Pairing	Bluetooth Low Energy with patented out-of-band pairing and AES-256 encryption.	Attested data channel from implant to node.	Cryptographic binding of implant to authorized devices; resists MITM.
On-implant Processing	Low-power amplification, digitization, and initial spike detection.	Trusted execution environment for generating commitments before wireless transmission.	Reduces attack surface on raw neural data.
Bidirectional Capability	Recording + stimulation (emerging).	Challenge-response protocols: node can issue neural “challenges” (subtle stimulation patterns) and verify responses.	Active liveness proof beyond passive recording.
Device Telemetry & State	Implant health metrics, thread integrity, power status,	Feed into Recursive Attestation (RFA) and substrate transition proofs.	Enables hardware-rooted continuity checks.

Primitive	Description	The Membrane Application	Security Properties
	firmware attestation.		

Caveats: - Current Neuralink devices are primarily medical (restoring function for paralysis, speech, etc.) with evolving long-term reliability (thread retraction observed in early trials).²⁰ - Full remote attestation capabilities are not publicly documented as of 2026; integration would require Neuralink API/SDK exposure or custom firmware extensions. - Neural data is extremely sensitive — all commitments must use strong privacy techniques (fuzzy extractors, homomorphic commitments, or zk-proofs over embeddings).²¹ - The N1 is a **single point of failure** inside the skull. If compromised at the hardware level, no software protocol can recover trust without surgical replacement.

Integration Points

1. Enhanced To Biological Anchor - During the genesis ceremony, derive the fuzzy biometric commitment from a combination of traditional biometrics **and** a Neuralink-derived neural signature (e.g., resting-state or task-evoked neural fingerprint). - Publish only the commitment + helper data to NOSTR. The raw high-dimensional neural template never leaves the secure channel. - **Limit:** If the neural template changes due to plasticity, injury, or device degradation, the fuzzy commitment must tolerate drift. This is harder than static biometric templates.

2. Neural-Enhanced Liveness Canary (Liveness-2 Circuit) Add as optional parallel domain: - **Neural Canary Event:** Participant performs a simple standardized mental or motor task (e.g., imagined movement, mental arithmetic, or specific imagery). - The N1 implant streams processed neural features (spike rates, LFP power in key frequency bands, population vector activity) to the The Membrane node. - The node generates a Merkle root over these features and includes it in the CP. - **ZK circuit augmentation:** The Liveness-1 circuit gains an additional column `neural_acc` that accumulates hash commitments of neural feature vectors. Boundary constraints verify the neural Merkle root against `pub_neural_root`.

What this actually adds: Higher-entropy, harder-to-simulate liveness evidence. The adversary must now simulate not just physics and optics, but the specific neural dynamics of a conscious brain performing a task.²²

What it does not add: Proof of consciousness. A sufficiently advanced adversary with a complete connectome model could still simulate the neural response. The cost is higher, not infinite.

3. Neural Challenge-Response (Active Liveness) When bidirectional stimulation matures: - The node issues a pseudorandom stimulation pattern (subtle, below perceptual threshold). - The N1 records the evoked neural response. - The node verifies that the response matches the expected signature for that identity and challenge. - This creates an **active** liveness proof: the entity being attested must have a live neural substrate capable of responding to perturbation, not just replaying pre-recorded data.

Caveat: Stimulation safety is medical-critical. Any challenge-response protocol must operate within clinically validated parameters. This is a medical ethics boundary, not just an engineering one.²³

4. Implant Telemetry as RFA Input - N1 firmware version, thread integrity metrics, and power state are included as public inputs to the node's RFA STARK. - If the implant reports thread retraction or firmware mismatch, the node fails closed pending surgical or remote re-anchoring. - **Caveat:** This telemetry comes from the implant itself. A compromised implant can lie about its own state. The telemetry is one input among many in cross-domain consensus, not a standalone truth.

Security Implications

Attack	Without Neuralink	With Neuralink
Replay pre-recorded sensor stream	Possible if video/IMU captured	Much harder: neural features are high-dimensional and task-specific
Deepfake video injection	Possible with generative AI	Does not bypass neural channel
Coerced channel use	Possible under duress	Active challenge-response may detect stress signatures (heuristic, not proof)
Hidden clone (never re-enters chain)	Undetectable	Still undetectable
Device compromise upstream	Possible if phone/wearable hacked	Requires compromising the implant itself or the BLE secure channel
Adversarial stimulation	N/A	Possible if attacker controls stimulation. Mitigated by medical-

Attack	Without Neuralink	With Neuralink
		safety bounds and cross-domain consensus.

Integration Roadmap

Phase	Milestone	Dependency
0	Documented Neuralink API/SDK for third-party attestation	Neuralink corporate decision
1	Fuzzy commitment prototype over neural embeddings	Research-only; no human trials without IRB
2	Liveness-2 circuit with neural column	Requires Winterfell AIR extension
3	Challenge-response protocol design	Medical safety validation; FDA or equivalent
4	Production integration	Long-term implant reliability proven (>5 years)

4.7 Agent-to-Agent Attestation Mode (Proposed Extension)

The Membrane’s core primitives — recursive Chain Proofs, Recursive Attestation (RFA), process continuity, and substrate-agnostic design — are directly applicable to **AI agent-to-agent (A2A) relations**. In a future with thousands of autonomous, delegating, and long-lived AI agents, the ability to cryptographically attest that an agent is running the expected process without undetected forking, rollback, or migration becomes critical. This mode extends **sovereignty** to pure-silicon entities.

For pure-silicon agents, the biological To anchor and human liveness provider are replaced by TEE substrate canaries. Agents can prove process integrity, reject unattested peers, and maintain verifiable delegation chains.

Why The Membrane Fits A2A Systems

Challenge	Solution
Silent forking / cloning	Recursive CP chain
Self-trust / Gödel problem	Recursive Attestation (RFA)
Mutual authentication	Cross-node consensus (2/3) + CP exchange
Delegation audit trails	Chained CPs across delegation hops

Challenge	Solution
Compromised agent containment	Fail-closed + substrate canary
Replay / stale state	Hybrid timestamp + continuity constraints
Operator capture	Self-hosted TEE + self-attestation

Core Adaptations for Pure Agent Mode

1. Agent Genesis (Silicon To) - A new agent is instantiated inside a TEE with a fresh keypair. - Initial CP references the code hash, model weights commitment (e.g., via Merkle root or zk-proof of weights), owner/operator identity (optional), and deployment configuration. - Published to NOSTR or a dedicated agent relay network.

2. Agent Liveness / Substrate Canary - Replaced by automated **substrate canary**: periodic TEE remote attestation proving: - Current code/memory hash matches expected state. - No unauthorized migration or checkpoint restore occurred. - Execution trace commitments (e.g., key nonces or decision logs) are consistent. - Cadence: every Δt (e.g., 1–60 minutes depending on agent risk level).

3. Agent-to-Agent Handshake Protocol When Agent A initiates contact with Agent B: 1. A sends its latest CP + RFA proof. 2. B verifies the chain and cross-node quorum. 3. B responds with its own CP. 4. Both parties may require additional attestations (e.g., task-specific state commitments) before proceeding. 5. High-stakes actions (financial transfers, code execution, data sharing) require fresh CPs within a short window.

Fail-closed rule: If either party cannot present a valid recent CP, the interaction is rejected or limited to read-only/low-trust mode. Interaction proceeds only when both parties present valid recent CPs.

4. Recursive Delegation Chains - When Agent A delegates a subtask to Agent B (who may further delegate), each hop appends a new CP referencing the parent task's CP. - The full delegation tree remains verifiable via recursive STARK composition. - This creates a tamper-evident audit trail for complex multi-agent workflows.

5. Integration with Existing A2A Standards - Compatible with DIDs / Verifiable Credentials (CP acts as a dynamic, high-assurance VC). - Can be layered over emerging protocols such as MCP, A2A, or agent communication frameworks. - NOSTR (or similar pub/sub mesh) serves as the lightweight messaging backbone for CP exchange.

Limitations in A2A Context

- **Performance overhead:** zk-STARK generation may be too slow for high-frequency, low-stakes interactions. Mitigation: cached recent proofs + lightweight substrate canaries for routine pings. See §10.1 for GPU acceleration path.
- **Social layer inapplicable:** Pure agents replace WoT with on-chain reputation, staking, or operator-signed endorsements. This introduces new trust assumptions.
- **Does not solve alignment:** The Membrane proves *process continuity and integrity*, not that the agent's goals are safe or aligned with the user/operator. Sovereignty is about self-determination, not benevolence.
- **TEE dependency:** Still limited by the security of underlying hardware enclaves.
- **Hidden clones:** As always, clones that never rejoin the attested chain remain undetectable.
- **Operator sovereignty conflict:** An agent may be “sovereign” in its attestation but still legally or practically owned by an operator. The protocol does not resolve ownership; it provides integrity evidence that can be used in governance or dispute resolution.

Security Implications Table

Attack	Without The Membrane	With Agent-E Mode
Forked agent impersonation	Easy (copy weights + state)	Requires forging entire CP chain + RFA
Stale / rolled-back agent	Common in long sessions	Prevented by continuity constraints
Malicious delegation chain	Hard to audit	Verifiable recursive audit trail
Mass agent compromise (supply-chain)	Catastrophic	Detectable via widespread CP failures and RFA mismatches
Operator hidden fork	Operator can silently fork	Detectable if fork must re-attest to interact with other agents
Unauthorized value extraction	Operator checkpoints, extracts, rolls back	Substrate canary detects checkpoint/restore

Implementation Roadmap Extension

Phase	Milestone (A2A)	Notes
0	Basic CP exchange over NOSTR	Pure silicon agents
1	Substrate canary + RFA for single agents	TEE-focused
2	Agent handshake protocol	Interoperability with DIDs
3	Delegation chain proofs	Multi-agent workflows
4	Full recursive composition at scale	High-value agent economies
5	Sovereign agent economies	Agents that own their own keys, host their own TEEs, and negotiate autonomously

Hybrid Human–AI Membrane

The biological T_0 anchor can gate **hybrid** stacks without merging human and agent sovereignty:

1. **Human holds T_0** and signs IACs for which channels may operate.
2. **Local or attested LLM** is a registered router channel (§4.2.2), not a silent default.
3. **Delegated agents** must chain CPs to `pub_parent_task_cp` referencing the human’s latest CP — preventing “my agent acted, but it wasn’t me” drift.
4. **Fail closed** applies uniformly: agent without valid parent CP + IAC cannot act on the subject’s behalf across the membrane.

This preserves compatibility with invasive BCI paths (§4.6) and pure-silicon A2A mode (§4.7) under one attestation framework.

4.8 Policy Enforcement vs Channel Continuity

Agent-governance stacks often focus on **policy enforcement** — proving an agent obeyed rules after the fact. The Membrane focuses on **channel continuity** — proving which read/write/routing paths were open and whether they matched prior commitments.

Dimension	Operator policy enforcement	The Membrane
Question answered	Did the agent obey policy?	Which channels were open and continuous?
Proof object	Conduct log / decision audit trail	Chain Proof of channel + anchor state
Trust locus	Operator, compliance, audit	Subject-local severance
Failure mode	Policy violation report	Membrane severance; channel kill

The two approaches are **composable**, not mutually exclusive: ILK-style decision logs (§Liveness-A) can support operator audit while the Membrane gives the **subject** a fail-closed boundary against routing and surveillance. Real-time A2A handshakes target sub-second verification and low-millisecond NOSTR publication overhead (§9.2, §10.1).

5. Formal Specification

Let **ProcessState(t)** be the cryptographic state at time t: memory commitments, execution context, nonce history. This is **not** consciousness. Let **F(t)** be the node attestation state. Let **L(t)** be the liveness canary event. Let **S(t)** be the social witness signature set.

Chain Proof CP(t):

```
CP(t) = zk-STARK_PROVE(
  witness = [L(t), S(t), CP(t-1), F(t), F(t-1)],
  public = [H(ProcessState(t)), H(F_code), timestamp, relay_root],
  constraints = [
    L(t) is within Δt of L(t-1),
    |S(t)| >= K,
    F(t) attests F(t-1) was valid per public rules,
    H(F_code) == audited_reference,
    H(ProcessState(t)) is consistent with H(ProcessState(t-1)) +
    delta
  ]
)
```

Verification: Any relay or node verifies CP(t) in $O(\log n)$ using the STARK proof and public inputs.^{[2425](#)}

What this actually proves: That a process with a specific history, running on attested infrastructure, generated a liveness signal and gathered social witness within a time bound. **Nothing MORE.**

6. Security Analysis

Primary Attack Matrix

Attack Vector	What the protocol actually does	What it cannot do
Fork / clone e-process	Detects <i>attested</i> forks by requiring CP chain continuity. Raises cost by requiring synchronized multi-domain attestation.	Cannot prove no hidden clone exists outside the attestation chain. Cannot prove subjective uniqueness.
VR epistemic capture	Ensures biological identity anchor remains active externally. Provides <i>external</i> evidence the target still has a grounded chain.	Cannot prevent a perfect simulation. Cannot prove the target is conscious inside VR.
Node self-compromise	Detected by cross-node consensus (2/3). Single node cannot forge historical chain of peers.	If 2/3 nodes compromised, consensus breaks. No escape from Byzantine bounds.
Replay / rollback	Prevented by hash-chain linking CP(t) to CP(t-1) and hybrid timestamping.	Requires rewriting chain from To. Does not prevent To compromise at genesis.
Sensor stream injection	Merkle root binds to specific sensor data.	If sensors or upstream pipeline compromised before Merkle root, injection is possible.
Coerced channel use	CP proves access to keys and registered devices.	Cannot distinguish willing use from coercion. Witnesses may detect

Attack Vector	What the protocol actually does	What it cannot do
LLM routing / sequestration	Router CP + IAC bind model id and context root; cross-domain disagreement fails closed.	anomalies (heuristic). Cannot detect routing through channels that never touch the membrane.
Cognitive identity drift	IAC TTL and session nonce raise cost of silent habituation to external inference.	Cannot prove subject still reasons endogenously when attestations remain valid.

Extended Attack Matrix (Neuralink + A2A)

Attack Vector	What the protocol actually does	What it cannot do
Neuralink adversarial stimulation	Active challenge-response may detect anomalous response patterns.	If attacker controls stimulation within medical bounds, may force “correct” response. Cross-domain consensus required for detection.
Neuralink firmware rollback	Firmware attestation hash checked in RFA.	If rollback occurs before attestation chain begins, undetectable.
Neuralink supply-chain implant	Detected if firmware hash mismatches audited reference at T_0 .	If backdoored firmware matches reference (colluding manufacturer), undetectable without destructive analysis.
Agent hidden fork by operator	Substrate canary + CP chain makes detectable if fork must re-attest. Delegation chain proofs require each hop to carry valid CP.	If fork never interacts with attested peers, undetectable. If compromised sub-agent never

Attack Vector	What the protocol actually does	What it cannot do
Agent delegation hijacking		rejoins chain, undetectable by upstream.
Agent value extraction (checkpoint/rollback)	Substrate canary detects unauthorized checkpoint/restore.	If extraction occurs between canary cycles, may be undetectable.

7. Recovery Mechanisms

This section defines explicit recovery paths for To loss, acknowledging that real systems must handle failure while preserving sovereignty.

7.1 Lost To Anchor (Hardware Failure)

Prevention: To key is sharded using Shamir’s Secret Sharing (3-of-5) across geographically distributed, trusted hardware tokens.²⁶

Recovery: 1. **Social recovery:** K witnesses from the original WoT perform a **re-anchor ceremony**. Each witness signs a new To’ event attesting that the same biological entity is present, using out-of-band verification (video call, in-person meeting, behavioral confirmation). 2. **Time-lock dead-man’s switch:** A pre-committed NOSTR event is published if no CP is received for $N \times \Delta t$. This event contains instructions for witness-initiated recovery or identity retirement. 3. **Biometric refresh:** A new fuzzy commitment is generated from the same biological template (with updated helper data). The old commitment is explicitly revoked on-chain.

Limitation: If the original WoT is also compromised or deceased, recovery fails. This is a social-layer failure mode, not a cryptographic one.

7.2 Compromised WoT Quorum

Detection: Sudden change in witness set, or valid CPs from witnesses who have not themselves published CPs recently.

Response: 1. **Quorum**

dilution: The target pre-commits a “witness refresh” event that can be activated with M-of-N signatures from a deeper recovery set (e.g., family members, legal trustees). 2. **Cooldown period:** New witnesses cannot sign

CPs until they have maintained valid CPs themselves for a minimum period (e.g., 30 days). 3. **Alert cascade:** If the WoT collapses, the dead-man's key triggers a hard sever and identity freeze until manual recovery.

7.3 Mass Witness Failure (Pandemic, War, Network Partition)

Graceful degradation: The protocol temporarily reduces K to a lower threshold K' (e.g., from 3 to 1) if a global alert condition is met (signed by 2/3 nodes + time-lock). This is an **explicit degradation**, not a silent failure. The target is notified that they are operating in reduced-trust mode.

7.4 Partial Compromise Handling

Domain Status	Other Domains Valid	Protocol Response
Liveness canary fails	WoT + node OK	Degraded mode: require additional witness confirmation within 6 hours, or fail closed.
WoT drops below K	Liveness + node OK	Degraded mode: accept CP with reduced K for one cycle, alert target, require WoT refresh within 48 hours.
Single node fails	Other 2 nodes + all domains OK	Continue operating. Failed node is marked for repair. RFA chain continues on remaining nodes.
2/3 nodes fail	Any domain status	Fail closed immediately. Dead-man's key becomes the only override.
Neuralink telemetry fails (if equipped)	All other domains OK	Degraded mode: fall back to Liveness-1 (traditional sensor) for this cycle. Alert target.
Neuralink firmware mismatch	All other domains OK	Fail closed. Surgical or remote re-anchoring required before resuming neural attestation.

Domain Status	Other Domains Valid	Protocol Response
Agent substrate canary fails	All other domains OK	Degraded mode: require fresh CP within shortened window, alert operator and peers.
Agent delegation chain break	Parent + other children OK	Isolate the broken branch, alert downstream agents, require re-attestation before resuming delegation.
Router session stale or model_id mismatch	Liveness + WoT + node OK	Sever LLM channel. Require fresh router CP + IAC or offline-only mode.
BCI + router active without cross-domain agreement	Either domain alone OK	Fail closed. Typical closed-loop sequestration path (EEG→LLM→output).
IAC expired or scope exceeded	Channels otherwise OK	Fail closed on affected channels until subject signs new IAC.

8. Economic and Game-Theoretic Analysis

8.1 Attack Economics

The protocol is designed to make sustained undetected compromise more expensive than the value of the target.

Attack	Estimated Cost	Protocol Defense
Forge one CP with $K=3$	Compromise 3 WoT humans + TEE + sensor pipeline	Recursive witness validity raises cost (each witness has own witnesses)
Maintain fork for 30 days	Continuous CP generation + social maintenance + no detection	Time-bounded liveness + out-of-band witness contact increases operational burden
Silent TEE compromise		

Attack	Estimated Cost	Protocol Defense
	Side-channel lab + physical access + firmware exploit	Cross-node consensus requires compromising 2/3 independent nodes
Perfect VR simulation	Full physics simulation + social graph simulation + NOSTR mesh simulation	“Physical expense” heuristic: simulation cost exceeds most attack budgets ²⁷
Neuralink simulation	Complete connectome model + real-time neural dynamics simulation	Cost orders of magnitude higher than physics simulation, but not infinite
Adversarial Neuralink stimulation	Medical device exploit + stimulation control	Constrained by medical safety bounds; cross-domain consensus detects anomalies
Agent fork for unauthorized delegation	Compromise operator + forge CP chain + maintain hidden clone	Detectable if fork must interact with attested peers; requires sustained operational cost
Mass agent supply-chain compromise	Backdoor all TEEs in agent fleet	Detectable via widespread CP failures and RFA mismatches

8.2 Incentive Structure

- **Witnesses:** No direct financial incentive. Reputation stake in the WoT (loss of trust if they attest falsely). Future: micro-rewards via Lightning Network for timely attestation.
- **Node operators:** Infrastructure cost only. Incentive is self-protection or community service. No tokenomics — avoids economic capture by AI or state actors.²⁸
- **Relays:** Standard NOSTR relay incentives (donation, subscription). No protocol-specific tokens.
- **Agents (A2A):** Operator reputation, task-completion rewards, or staking slashing. The protocol itself does not enforce incentives; it provides the attestation layer on which incentive mechanisms can be built.

8.3 Coercion Resistance

The protocol cannot prevent coercion. It can: - Raise the cost by requiring *continuous* coercion across Δt (not one-time extraction). - Enable detection via behavioral anomaly (witnesses who know the target notice deviations). - Provide a duress mode: a specific CP variant that signals coercion to the WoT without alerting the coercer. - **Neuralink extension:** Active challenge-

response may detect stress/anxiety signatures (elevated amygdala activity, altered prefrontal coherence), but this is heuristic, not proof. - **Agent context:** An operator may coerce an agent by forcing it to run unauthorized code. The substrate canary detects code hash mismatch, but if the operator controls the TEE, they may suppress the canary. This is why self-hosted TEEs (not cloud-hosted) are critical for agent sovereignty.

9. Minimum Viable Prototype (MVP)

9.1 Definition

Phase 0 demonstrates a **self-hosted cognitive firewall** without requiring implants or VR.

MVP components: 1. Self-hosted NOSTR relay 2. TEE node (single AMD SEV-SNP or Intel TDX instance) 3. **Intent gate:** Subject-signed IAC before router or BCI decode (§4.2.1) 4. **AI router gate:** Local or attested inference session with context-hash commitment in CP (§4.2.2) 5. **Optional BCI/sensor path:** Smartphone or EEG feature Merkle root (Liveness-1/2 circuit) 6. Witness set: K=2–3 trusted humans (WoT) 7. zk-STARK circuit: channel commitments + timestamp + witness count 8. Dead-man’s key: hardware token for emergency severance

The MVP proves a subject can maintain membrane-gated channels without routing cognition through unattested corporate endpoints by default.

9.2 Success Metrics

Metric	Target	Evaluation Method
Proof generation time	< 5 minutes on consumer hardware (desktop)	Benchmark Winterfell prover with 60-row trace
Proof size	< 50 KB	Measure STARK output bytes
Verification time	< 100 ms	Benchmark verifier on standard CPU
Witness confirmation latency	< 4 hours	Measure NOSTR gossip propagation + human response time
False positive rate (fail)	< 1% per month	Log all false alerts and root-cause

Metric	Target	Evaluation Method
closed without attack)		
Operational continuity (target)	30 days without manual intervention	Time-to-first-failure test
Neuralink extension proof size (future)	< 75 KB	Estimate additional column overhead
A2A handshake latency (future)	< 2 seconds	Cached CP + lightweight substrate canary
Proof verification latency (A2A target)	< 250 ms	Benchmark verifier on standard CPU; required for real-time agent handshakes
NOSTR publication overhead (target)	< 10 ms	Measure event publish latency to self-hosted relay

9.3 Dependency Separation

Immediately Actionable	Future-Dependent
NOSTR relay + OpenClaw	BCI integration
Smartphone sensor pipeline	Full VR runtime attestation
Winterfell STARK circuit	Multi-party recursive composition at scale
Single TEE node	Cross-vendor TEE consensus (AMD + Intel + ARM)
K=2 WoT	Large-scale WoT with recursive witness validation
Liveness-1 (traditional sensors)	Liveness-2 (Neuralink neural signals)
Basic CP exchange	Full A2A delegation chains
CPU proving	GPU-accelerated proving (§10.1)

10. Implementation Roadmap

Phase	Deliverable	Stack	Blockers
0	NOSTR attestation bus + OpenClaw automation	Rust/TypeScript, NOSTR relays ²⁹	Relay censorship, Sybil relays
1	Liveness canary zk-STARK circuit	Winterfell/Cairo, SHA-256/Rescue ³⁰	STARK proof generation latency on mobile/embedded
2	Self-hosted TEE node	AMD SEV-SNP, Linux KVM ³¹	Supply-chain trust, side-channel leakage
3	Social graph witness protocol	NOSTR WoT, gossip model ³²	Social layer compromise, witness coercion
4	Substrate transition gate	VR runtime integration ³³	VR runtime itself is a massive TCB
5	Recursive node attestation	Multi-party STARK composition ³⁴	2/3 compromise threshold is hard limit
6	Neuralink extension (optional)	N1 API + Winterfell AIR extension ³⁵	Medical device maturity, corporate API availability, FDA validation
7	Agent-to-Agent mode	Pure-silicon CP + handshake protocol + delegation chains	Performance overhead, incentive design, TEE availability for agents
8	GPU acceleration (optional)	CUDA kernels + Winterfell backend	NVIDIA dependency, driver-level trust, mobile power constraints

10.1 Hardware Acceleration: CUDA and GPU Support (Proposed Extension)

Proof generation remains one of the primary practical blockers for frequent attestation (short Δt) and complex circuits (Liveness-2 neural features, recursive RFA, A2A delegation chains). GPU acceleration, particularly via CUDA on NVIDIA hardware, offers a high-impact optimization path toward sub-second verification and low-latency NOSTR publication.

Relevance to The Membrane

- **Target workloads:** Polynomial arithmetic (FRI, DEEP), Merkle tree hashing, constraint evaluation, and recursive STARK composition.
- **Expected impact:** 5–50x+ speedup on key kernels based on related STARK GPU projects (e.g., miniSTARK, ICICLE-inspired work). This enables:
 - Mobile/consumer-grade proving for daily liveness.
 - Higher-frequency substrate canaries for agents.
 - Larger witness sets in social consensus.

Implementation Approach

- **Optional CUDA backend** for Winterfell-compatible prover: Accelerate field operations, NTT/FFT, and hash chains using CUDA kernels (or Rust CUDA abstractions).
- **Fallback path:** Pure CPU (multi-threaded) remains fully supported. Detection of CUDA-capable hardware enables accelerated mode automatically.
- **Cross-platform consideration:** Prioritize CUDA for high-performance nodes, while exploring WebGPU / Vulkan SPIR-V (via rust-gpu) for broader device compatibility (including future mobile GPUs).
- **Sovereignty note:** Users retain full control — acceleration runs locally on self-hosted hardware. No cloud GPU dependency required.

Caveats

- Does **not** change security or privacy properties.
- Increases binary size and introduces driver-level trust (mitigated by open-source kernels and verifiable computation).
- NVIDIA-specific initially; long-term goal is vendor-agnostic GPU support.
- Power/thermal constraints on embedded or mobile devices may limit always-on acceleration.

Roadmap Integration

Phase	Milestone	Benefit
1+	CUDA-accelerated polynomial & hash kernels	Faster MVP proving on desktops/servers
2+	Integrated Winterfell GPU backend (optional feature)	Mobile-friendly liveness proofs
3+	GPU-optimized recursive composition	Practical high-frequency A2A & Neuralink use

11. Open Source, Audit, and Interoperability

11.1 Licensing

Reference implementation: **GPL-3.0** or **AGPL-3.0**. The protocol must remain open and auditable. Patent license: explicit defensive patent pledge for all contributors.

11.2 Audit Plan

Stage	Auditor Type	Scope
Pre-alpha	Community audit (open source)	NOSTR event schema, circuit logic
Alpha	Independent security firm	TEE integration, side-channel analysis
Beta	Academic cryptographers	zk-STARK soundness, recursive composition
Production	Multi-party audit consortium	Full stack, social recovery, coercion resistance, Neuralink extension safety, A2A delegation integrity

11.3 Interoperability

- **NOSTR:** Native. Uses standard kinds 31990/31991. Compatible with existing clients as read-only viewers.
 - **TEE ecosystems:** Targets AMD SEV-SNP and Intel TDX via standard attestation APIs. ARM CCA as future target.
 - **Identity bridges:** Optional DID:web or DID:ion mapping for compatibility with W3C decentralized identity systems. No dependency.
 - **ZK ecosystems:** STARK proofs verifiable on StarkNet or Cairo VM for on-chain anchoring if desired. Not required.
 - **Neuralink:** Requires proprietary API/SDK. If unavailable, the protocol falls back to Liveness-1 without degradation.
 - **A2A frameworks:** Can be layered over MCP, A2A, or custom agent communication protocols. CP exchange is transport-agnostic.
-

12. Ethical Considerations

12.1 Digital Minds and Agent Exit Rights

If e-processes or pure agents achieve legal or moral personhood, the membrane must not become a tool of digital imprisonment. The dead-man's key and fail-closed behavior must preserve the right to exit — including the right to sever the attestation chain without operator permission.

12.2 Coercion and Autonomy

The membrane's channel requirements could be weaponized: a captor forces the subject to maintain attested sessions. The duress signal and dead-man's key are partial mitigations, not solutions.

Neuralink extension adds new coercion vectors: - Forced implantation or replacement. - Adversarial stimulation to induce compliance. - Neural data extraction under duress. - Mitigation: explicit informed consent requirements, right to refuse implantation, and medical ethics oversight.

Agent coercion: An operator may force an agent to run unauthorized code or delegate to compromised peers. The substrate canary detects code hash mismatch, but if the operator controls the TEE, suppression is possible. Self-hosted TEEs are the primary mitigation.

12.3 Right to Be Forgotten

To anchors are permanent by design. A revocation mechanism must exist that allows an identity to be **explicitly retired** (not recycled) with a final CP signed by the original WoT. This is irreversible and prevents zombie identity attacks.

13. Known Open Problems

These are explicitly unresolved and invite community contribution. **Top-priority problems are marked with ★.**

1. Formal verification of the recursive composition. No complete machine-checked proof exists for the multi-party STARK RFA chain. This is the highest-priority open problem for protocol credibility.

2. High-frequency A2A attestation overhead. zk-STARK generation may be too slow for real-time agent handshakes. Lightweight alternatives (cached proofs + Merkle state diffs) or GPU acceleration (§10.1) need research.

3. Long-term neural template drift and fuzzy extractor stability. Neural patterns change over time due to plasticity, aging, injury, and device degradation. Fuzzy commitments that tolerate this drift while maintaining security are an open research problem.^{[3637](#)}

- 1. Mobile STARK generation.** Current provers are too slow for smartphone-native proof generation. Requires circuit optimization or delegated proving with privacy preservation.
- 2. TEE side-channel resistance.** No TEE currently offers provable resistance to all known side-channel attacks. Hardware advances needed.
- 3. Coercion-proof authentication.** All liveness signals are vulnerable to coercion. A coercion-proof primitive would require fundamental advances in behavioral biometrics or neurochemistry.
- 4. Social graph Sybil resistance at scale.** The WoT model works for small K. Scaling to thousands of witnesses while maintaining recursive validity is an open graph-theoretic problem.
- 5. Quantum-resistant hash functions for STARKs.** While hash-based cryptography is post-quantum, specific hash functions used in STARKs (e.g., Rescue, Poseidon) require continued cryptanalysis.

6. **Economic sustainability without tokens.** The no-token incentive model relies on altruism and self-interest. Long-term sustainability of relay and witness networks is unproven.
7. **Agent sovereignty vs. operator ownership.** The protocol proves agent integrity, but does not resolve legal or economic ownership. If an operator owns the hardware, can they claim to own the agent's attestations? This is a governance and legal problem, not a cryptographic one.

★ 11. **Cognitive identity drift detection.** Distinct from neural template drift (problem 3): the attestation chain remains valid while the subject's effective reasoning migrates to external LLM/BCI closed loops. Detecting routing substitution and epistemic capture **without reading thought content** — using IAC scope, router CP consistency, and cross-domain disagreement only — is an open problem tied to §0.2 and §4.2.³⁸

Part 2: zk-STARK Circuit Architecture

Circuit Comparison

Property	Liveness-1	Liveness-2	Liveness-A
Target entity	Biological human	Biological human + BCI	Pure silicon agent
Columns	8	9 (+ neural_acc)	7 (+ ilk_acc)
Hash accumulator	sensor_acc	sensor_acc + neural_acc	trace_acc + ilk_acc
Witness layer	Social graph (≥K humans)	Social graph (≥K humans)	Optional operator endorsement
Liveness source	IMU / video / audio	IMU / video / audio + neural features	Execution trace commitments
TEE attestation	node only	node + implant telemetry	Agent enclave attestation
Recursive input	pub_node_proof_hash	pub_node_proof_hash	pub_node_proof_hash + pub_parent_task_cp
Proof size	~25 KB	~35 KB	~25 KB

Property	Liveness-1	Liveness-2	Liveness-A
Use case	Channel integrity (human)	BCI-augmented channel integrity	Agent substrate canary
Dependencies	Smartphone / wearable	BCI device (future)	Self-hosted TEE

Liveness-1 Circuit (Traditional Sensors)

Overview

This circuit proves that **some entity** with access to the biometric key and sensor devices generated a liveness signal within a time window, linked to an identity chain, and gathered social attestation. It explicitly does **not** prove consciousness, free will, or absence of coercion.

Property	Value
Circuit Name	Liveness-1
Proof System	zk-STARK (hash-based, transparent setup) ³⁹⁴⁰
Hash Function	Poseidon2 (AIR efficiency) + SHA-256 (NIST PQC alignment) ⁴¹
Field	64-bit Goldilocks
Target	Winterfell (Rust) or Cairo VM ⁴²

Public Inputs

Register	Description
pub_identity	NOSTR npub (compressed)
pub_prev_cp_hash	Hash of CP(t-1)
pub_timestamp	Hybrid Unix timestamp
pub_merkle_root	Merkle root of sensor chunk hashes
pub_witness_count	Valid witness signatures ($\geq K$)
pub_witness_root	Merkle root of witness npubs (anonymized)
pub_node_hash	Hash of node code + TEE attestation

Private Inputs

Register	Description
priv_sensor_data	IMU vectors, video frame hashes, audio fingerprints
priv_biometric_key	Derived key from fuzzy commitment of To anchor ⁴³
priv_witness_sigs	Witness signatures on NOSTR event
priv_location_entropy	Optional noise hash (omittable)

Execution Trace (AIR)

8 columns, N rows (scales with witness count).

Col	Name	Role
0	clk	Clock counter
1	sensor_acc	Accumulator for sensor data hash chain
2	witness_acc	Accumulator for verified witness signatures
3	time_delta	Decrementing: $\Delta t - clk$
4	continuity	Boolean: 1 if sensor_acc matches pub_merkle_root at final row
5	sig_valid	Boolean: 1 if witness signature verifies
6	threshold	Boolean: 1 if witness_acc $\geq K$ at final row
7	prev_link	Hash of pub_prev_cp_hash pub_identity pub_timestamp

Transition Constraints

C1 — Sensor Chain Integrity:

`sensor_acc' = Hash(sensor_acc || priv_sensor_data[row])`

*Enforces Merkle root is built from provided sensor data. **Does not prove sensors were honest.***⁴⁴

C2 – Witness Accumulator:

```
witness_acc' = witness_acc + sig_valid
```

C3 – Time Bound:

```
time_delta' = time_delta - 1  
time_delta >= 0
```

C4 – Signature Verification:

```
sig_valid = VerifyEd25519(witness_sig, witness_npub, event_hash)
```

C5 – Continuity Link:

```
prev_link = Hash(pub_prev_cp_hash || pub_identity ||  
pub_timestamp)
```

Boundary Constraints

First row: sensor_acc = 0, witness_acc = 0, time_delta = Δt , continuity = 0

Last row: sensor_acc == pub_merkle_root, witness_acc >= K, time_delta >= 0, continuity == 1

Recursive Composition (RFA)

Additional public input:

Register	Description
pub_node_proof_hash	Hash of node's own STARK proof from T-1

Last-row constraint:

```
prev_link = Hash(pub_prev_cp_hash || pub_node_proof_hash ||  
pub_identity)
```

This creates a **proof chain**. An adversary compromising the node at T must forge the entire historical chain of node attestations across $\geq 2/3$ independent nodes.⁴⁵

Liveness-2 Neural Extension (Optional)

Overview

The Liveness-2 circuit extends Liveness-1 by adding a **neural feature accumulator** column. It is designed for BCI-equipped users (e.g., Neuralink N1) and operates as an optional parallel domain. The protocol accepts either Liveness-1 or Liveness-2 signals, or both for higher confidence.

Additional Public Inputs

Register	Description
pub_neural_root	Merkle root of neural feature chunk hashes
pub_neural_task_id	Identifier of the standardized neural task performed (e.g., “imagined_fist_5s”)

Additional Private Inputs

Register	Description
priv_neural_data	Processed neural features: spike rates, LFP power bands, population vectors

Modified Execution Trace (AIR)

The Liveness-2 trace has **9 columns** (adding column 8):

Col	Name	Role
0-7	<i>(same as Liveness-1)</i>	<i>(same as Liveness-1)</i>
8	neural_acc	Accumulator for neural feature hash chain

Additional Transition Constraints

C6 — Neural Chain Integrity:

```
neural_acc' = Hash(neural_acc || priv_neural_data[row])
```

*Enforces neural Merkle root is built from provided neural features. **Does not prove the neural data came from a conscious brain.***⁴⁶

Modified Boundary Constraints

First row: `neural_acc = 0` (added to existing first-row constraints)

Last row: `neural_acc == pub_neural_root` (added to existing last-row constraints)

Security Properties

- Proves that an entity with the biometric key generated a neural-feature Merkle root within Δt for a specific task.
 - Does **not** prove the neural data was not synthesized by a connectome model.
 - Does **not** prove the entity was conscious during the task.
 - Does **not** prove the N1 implant was not backdoored at manufacture.
 - Active challenge-response (future) adds an additional constraint: the neural response must match a stimulation-locked pattern, raising simulation cost.⁴⁷
-

A2A Agent Circuit (Liveness-A)

Overview

The Liveness-A circuit is a simplified variant for pure-silicon agents. It removes the biological/social witness layer and replaces it with substrate attestation, an Immutable Logging Kernel (ILK) for complete decision audit trails, and optional operator endorsement.

Public Inputs

Register	Description
<code>pub_agent_identity</code>	Agent's NOSTR npub or DID
<code>pub_prev_cp_hash</code>	Hash of CP(t-1)
<code>pub_timestamp</code>	Hybrid Unix timestamp
<code>pub_code_hash</code>	Hash of agent code / model weights commitment
<code>pub_tee_attestation</code>	TEE remote attestation report
<code>pub_node_hash</code>	Hash of node code + TEE attestation
<code>pub_parent_task_cp</code>	Hash of parent task CP (for delegation chains)
<code>pub_ilk_root</code>	

Register	Description
	Merkle root of Immutable Logging Kernel decision log

Private Inputs

Register	Description
priv_execution_trace	Commitments to key nonces, decision logs, or state transitions
priv_decision_log	Immutable Logging Kernel (ILK) — Merkle root of all agent decisions since last CP
priv_operator_sig	Optional operator endorsement signature

Execution Trace (AIR)

6 columns, N rows.

Col	Name	Role
0	clk	Clock counter
1	trace_acc	Accumulator for execution trace commitments
2	ilk_acc	Accumulator for Immutable Logging Kernel decision hashes
3	time_delta	Decrementing: $\Delta t - clk$
4	continuity	Boolean: 1 if trace_acc matches pub_code_hash at final row
5	tee_valid	Boolean: 1 if TEE attestation verifies
6	prev_link	Hash of pub_prev_cp_hash pub_agent_identity pub_timestamp

Transition Constraints

C1 — Execution Trace Integrity:

```
trace_acc' = Hash(trace_acc || priv_execution_trace[row])
```

C2 – Immutable Logging Kernel (ILK):

```
ilk_acc' = Hash(ilk_acc || priv_decision_log[row])
```

Enforces that every agent decision since the last CP is committed into the attestation chain. This creates a complete behavioral audit trail that proves not just code integrity, but decision continuity.

C3 – TEE Attestation Verification:

```
tee_valid = VerifyTEEAattestation(pub_tee_attestation,  
pub_code_hash)
```

C4 – Continuity Link:

```
prev_link = Hash(pub_prev_cp_hash || pub_agent_identity ||  
pub_timestamp)
```

Boundary Constraints

First row: trace_acc = 0, ilk_acc = 0``, time_delta = Δt , continuity = 0, tee_valid = 0`

Last row: trace_acc == pub_code_hash, ilk_acc == pub_ilk_root, time_delta >= 0, continuity == 1, tee_valid == 1

Sovereignty Note

The Liveness-A circuit proves **process integrity**, not **goal alignment**. A sovereign agent can prove it is running the code it claims to be running, but this does not prove the code is safe or benevolent. Sovereignty is about self-determination and integrity, not about trustworthiness.

Integration with NOSTR + OpenClaw

NOSTR bus: - CP events: kind 31990, tags ["e", <prev_event_id>], ["p", <firewall_npub>], ["k", "the-membrane-liveness"] - Neural-enhanced CP events: kind 31990, tags ["k", "the-membrane-liveness-neural"] (optional) - BCI channel CP events: kind 31990, tags ["k", "the-membrane-bci"] (optional) - Router session CP events: kind 31990, tags ["k", "the-membrane-router"], plus ["model", <model_id>], ["ctx", <context_merkle_root>] (optional) - IAC events: kind 31990, tags ["k", "the-membrane-iac"], content = signed scope JSON (optional) - Agent CP

events: kind 31990, tags ["k", "the-membrane-agent"] (optional) - Relays: self-hosted + 2+ public for redundancy. **Caveat:** relays can censor or partition.⁴⁸

OpenClaw automation: - Trigger: IMU threshold or cron. - Action: Compile sensor data → generate STARK → publish NOSTR. - Fallback: Proof failure or missing witnesses → post kind 31991 “degraded continuity” alert → fail closed.

Citation Tiers (Explicit Epistemic Weight)

Tier	Type	Citations
T1: Formal	Peer-reviewed or established cryptographic standards	⁴⁹ Neulinger & Sparer (Springer), ⁵⁰ Juels & Wattenberg (ACM CCS), ⁵¹ Ben-Sasson et al. (STARKs), ⁵² NIST PQC, ⁵³ Shamir (ACM)
T2: Implementation	Production systems, protocols, hardware docs	⁵⁴ AMD SEV-SNP / Intel TDX, ⁵⁵ Winterfell/Cairo, ⁵⁶ NOSTR protocol, ⁵⁷ Neuralink Corp (medical device docs)
T3: Conceptual	Speculative architecture, sci-fi framing, heuristic arguments	⁵⁸ Khan preprint (cognitive firewall concept), ⁵⁹ Banks (Surface Detail as inverse reference), ⁶⁰ embodied cognition heuristic

No citation in this document is presented as stronger than its tier.

Appendix A: Version History

Version	Date	Notes
v0.9.5	2026-06-14	Remove third-party competitor citations (policy vs continuity §4.8 retained)
v0.9.4	2026-06-14	

Version	Date	Notes
		Independent research / no-affiliation disclaimer (title block, §0.5, contribute)
v0.9.3	2026-06-14	IAC layer (§4.2.1), router session CPs (§4.2.2), epistemic capture (§0.2), substrate parity, policy vs continuity (§4.8), hybrid human–AI membrane, cognitive identity drift open problem
v0.9.2	2026-06-14	Appendix B: open research & Phase 0 stack; published to GitHub
v0.9.1	2026-06-14	Remove cultured tissue (§4.8) and Related Work (§14)
v0.9	2026-06-14	Refocus on cognitive firewall; remove DojoPop conflation
v0.8	2026-06-14	Strip AI slop; pluggable channels; rebrand from GORGONEUM
v0.6	2026-05-12	Related work, GPU section, A2A and Neuralink extensions
v0.5	2026-05-12	Sovereignty framing and agent attestation mode

Appendix B: Open Research & Prototype Stack

Neuralink: what is public

Neuralink has **no open SDK** and no third-party attestation API for the clinical N1 stack. Useful public material:

Resource	URL	Use for Membrane
ASIC / dev platform paper	bioRxiv 703801	Streaming model, electrode counts, UDP multicast — not consumer app stack
Patient / device overview	neuralink.com/patient-information	Threat model for BLE pairing, firmware updates, phone app as TCB
PRIME study	ClinicalTrials NCT06429735	Regulatory context only

Treat Neuralink as a **future Liveness-2 target**, not a build dependency today.

Open acquisition & decode (Phase 0 hardware)

Project	License	Role
BrainFlow	MIT	Unified SDK for OpenBCI, Muse, etc. — channel registration + feature extraction
OpenBCI	Open hardware + OSS tooling	Lab bench; Galea adds EEG/EMG/PPG/EDA in XR
Lab Streaming Layer	BSD	Time-synced multimodal bus — use as membrane data plane

Project	License	Role
MetaBCI	GPL-2.0	Full decode pipeline; study decoder drift when models retrain
PyNoetic	GPL-3.0	Real-time EEG → classifier pipelines

Invasive / bench research (not consumer implants)

Project	Notes
Iris-128	Open 128-ch headstage (CERN-OHL)
OpenMEA	Benchtop closed-loop electrophysiology

Security & firewall research (closest to Membrane thesis)

Project	Type	Membrane mapping
iba-neural-guard	Open repo	Decoded signal ≠ authorization. Signed intent certificate before any BCI action; blocks capability drift
Argus / LibArgus	Paper (no public repo)	OS information flow control for wearable BCIs; 300+ vulns on Muse/NeuroSky/OpenBCI stacks
NSP — Neural Sensory Protocol	Spec draft	PQC wire protocol for neural frames; constant-rate BLE against timing side channels
NeuroZKP	Project site	ZK proofs over neural signals — prove authorization without raw spike disclosure

iba-neural-guard is the closest open articulation of the policy layer Membrane needs: the thought is not the authorization; the signed certificate is.


LLM [?] brain coupling (adversary models)

Active research on **closed-loop brain** → **LLM routing** — where identity drift and sequestration appear first:

Project	URL	Why it matters
NeuroLM	github.com/935963004/NeuroLM	EEG treated as “foreign language” inside an LLM — native cognition absorbed into model space
SYNAPTICON	github.com/AlbertBarqueDuran/SYNAPTICON	EEG → text → LLM → output closed loop
Brain-LLM Interface	arXiv:2603.16897	EEG gates LLM refinement at inference time

Membrane term	Meaning in this stack
Identity drift	Decoder or LLM updates, session fork, model swap without new Chain Proof
Sequestration	Cognition lives in external inference; endogenous loop bypassed or unreadable
Firewall job	Gate which channels may couple; sever on missing/stale CP

Suggested Phase 0 architecture



Build order

1. **BrainFlow + LSL** — stream EEG features; never publish raw traces to NOSTR.

2. **Channel registry** — YAML list of permitted paths (BCI app, local LLM port, forbidden cloud URLs).
3. **Merkle commitment circuit** — Winterfell Liveness-1 over feature vectors + timestamp (see whitepaper Part 2).
4. **Intent gate** — adapt [iba-neural-guard](#) pattern: no signed scope → no decode→action mapping.
5. **NOSTR CP bus** — kind 31990, tags #the-membrane, hash chain to prior event.
6. **WoT** — K=2 human witnesses sign CP validity out-of-band (Signal/video), not live ritual watching.

What Phase 0 proves

- Subject can run membrane-gated channels **without default cloud LLM routing**.
- Unauthorized channel open → detectable in attestation chain (if channel touches the bus).
- Subject can **sever** all links (dead-man's key / kill switch).

What Phase 0 does not prove

- Covert passive surveillance with no attested channel.
 - Thought content or consciousness.
 - Neuralink-class invasive implant integration.
-

References to add to implementation issues

- Khan, E. (2026). *Brain Hacking: AI for Safeguarding Against Dangerous AI*. [Preprints.org](#) — cognitive firewall framing (T3).
- Li et al. BCI cybersecurity survey — [arXiv:2007.09466](#) — Neuralink BLE threat model.
- Argus wearable BCI security — [arXiv:2201.07711](#).

How to Contribute

This is **community-facing independent research**. Contributors are welcome; contribution does not create an organizational entity or imply affiliation with the author's employers or clients.

Implement the MVP (§9) using the stacks in [Appendix B](#). Publish benchmarks via GitHub issues. Security disclosures: notify the author privately, 90-day coordinated disclosure. Reference implementation: GPL-3.0/AGPL-3.0.

Footnotes

Tags: [#the-membrane](#) [#whitepaper](#) [#zk-stark](#) [#identity-attestation](#) [#process-continuity](#) [#sovereignty](#) [#nostr](#) [#agent-attestation](#)

1. NeuroLM, SYNAPTICON, Brain-LLM Interface — see [Appendix B](#). **Conceptual / active research (T3)**. Closed-loop EEG→LLM pipelines illustrate sequestration and identity drift threats; not production security claims.[\[?\]](#)
2. Chalmers, D. J. (1995). *Facing Up to the Problem of Consciousness*. Journal of Consciousness Studies. — **Philosophical baseline (T3)**. The “hard problem” remains unsolved. No cryptographic protocol can bridge the explanatory gap.[\[?\]](#)
3. Gödel, K. (1931). *Über formal unentscheidbare Sätze der Principia Mathematica und verwandter Systeme I*. — **Formal limit (T1)**. Any self-referential proof system faces consistency limits. Cross-domain consensus is the practical workaround, not a solution.[\[?\]](#)
4. Heuristic argument. — **Conceptual / speculative (T3)**. The claim that physical-sensor signals are “expensive to simulate” is a cost argument, not a cryptographic proof. Vulnerable to replay, injection, and coercion.[\[?\]](#)
5. AMD. *SEV-SNP: Strengthening VM Isolation with Integrity Protection and More*. White paper. Intel. *Trust Domain Extensions (TDX) Architecture Specification*. — **Implementation / hardware docs (T2)**. TEEs provide best-available isolation but are vulnerable to side-channels and supply-chain attacks.[\[?\]](#)
6. National Institute of Standards and Technology (NIST). *SP 800-207: Zero Trust Architecture*. <https://doi.org/10.6028/NIST.SP.800-207> — **Standards (T1)**. Vendor-neutral reference for Zero Trust Architecture. The Membrane implements these principles via cryptographic consensus rather than network segmentation.[\[?\]](#)
7. Neulinger, J. & Sparer, K. (2025). *Fostering AI alignment through blockchain, proof of personhood and zero knowledge proofs*. Cluster Computing. <https://link.springer.com/article/10.1007/>

s10586-025-05729-8 — **Formal / peer-reviewed (T1)**. Establishes PoP consensus, zk-STARK identity verification, AI Shield model. Core engineering reference.[?]

8. Neuralink Corporation. *Patient Information: N1 Implant and R1 Robot*. <https://neuralink.com/patient-information/> — **Implementation / medical device (T2)**. As of 2026, Neuralink devices are investigational medical devices with demonstrated thread retraction in early trials. No public remote attestation API exists. All security claims are prospective.[?]
9. Neuralink Corporation. *Patient Information: N1 Implant and R1 Robot*. <https://neuralink.com/patient-information/> — **Implementation / medical device (T2)**. As of 2026, Neuralink devices are investigational medical devices with demonstrated thread retraction in early trials. No public remote attestation API exists. All security claims are prospective.[?]
10. Chalmers, D. J. (1995). *Facing Up to the Problem of Consciousness*. Journal of Consciousness Studies. — **Philosophical baseline (T3)**. The “hard problem” remains unsolved. No cryptographic protocol can bridge the explanatory gap.[?]
11. Gödel, K. (1931). *Über formal unentscheidbare Sätze der Principia Mathematica und verwandter Systeme I*. — **Formal limit (T1)**. Any self-referential proof system faces consistency limits. Cross-domain consensus is the practical workaround, not a solution.[?]
12. Juels, A. & Wattenberg, M. (1999). *A Fuzzy Commitment Scheme*. ACM CCS. — **Formal (T1)**. Allows key derivation from noisy biometrics without storing raw templates.[?]
13. Grokipaedia. *iba-neural-guard*. <https://github.com/Grokipaedia/iba-neural-guard> — **Implementation reference (T2)**. Intent certificate pattern: decoded neural signal $\neq$ authorization. Adapted for Membrane IAC layer (§4.2.1).[?]
14. Zimmermann, P. (1995). *PGP User's Guide*. — **Implementation reference (T2)**. Web-of-trust model adapted for biological witness attestation.[?]
15. Zimmermann, P. (1995). *PGP User's Guide*. — **Implementation reference (T2)**. Web-of-trust model adapted for biological witness attestation.[?]
16. TPM / TEE remote attestation literature. — **Implementation analogy (T2)**. Substrate canaries for silicon processes mirror TPM attestation chains but apply to runtime state rather than boot state.[?]

17. Banks, I. M. (2010). *Surface Detail*. Orbit. — **Fiction / conceptual inverse (T3)**. Depicts virtual Hells without integrity verification or exit. Used as architectural counter-example, not evidence. [?]
18. AMD. *SEV-SNP: Strengthening VM Isolation with Integrity Protection and More*. White paper. Intel. *Trust Domain Extensions (TDX) Architecture Specification*. — **Implementation / hardware docs (T2)**. TEEs provide best-available isolation but are vulnerable to side-channels and supply-chain attacks. [?]
19. Lamport, L., Shostak, R., & Pease, M. (1982). *The Byzantine Generals Problem*. ACM TOPLAS. — **Formal (T1)**. Establishes the 2/3 consensus bound. No recursive mechanism can exceed this threshold without additional trust assumptions. [?]
20. Neuralink Corporation. *Patient Information: N1 Implant and R1 Robot*. <https://neuralink.com/patient-information/> — **Implementation / medical device (T2)**. As of 2026, Neuralink devices are investigational medical devices with demonstrated thread retraction in early trials. No public remote attestation API exists. All security claims are prospective. [?]
21. Juels, A. & Wattenberg, M. (1999). *A Fuzzy Commitment Scheme*. ACM CCS. — **Formal (T1)**. Allows key derivation from noisy biometrics without storing raw templates. [?]
22. Heuristic argument. — **Conceptual / speculative (T3)**. The claim that physical-sensor signals are “expensive to simulate” is a cost argument, not a cryptographic proof. Vulnerable to replay, injection, and coercion. [?]
23. Neuralink Corporation. *Patient Information: N1 Implant and R1 Robot*. <https://neuralink.com/patient-information/> — **Implementation / medical device (T2)**. As of 2026, Neuralink devices are investigational medical devices with demonstrated thread retraction in early trials. No public remote attestation API exists. All security claims are prospective. [?]
24. Neulinger, J. & Sparer, K. (2025). *Fostering AI alignment through blockchain, proof of personhood and zero knowledge proofs*. Cluster Computing. <https://link.springer.com/article/10.1007/s10586-025-05729-8> — **Formal / peer-reviewed (T1)**. Establishes PoP consensus, zk-STARK identity verification, AI Shield model. Core engineering reference. [?]

25. Ben-Sasson, E. et al. (2018). *Scalable, transparent, and post-quantum secure computational integrity*. IACR ePrint. — **Formal (T1)**. Foundational STARK paper. Winterfell (Rust) and Cairo are production implementations.^[?]
26. Shamir, A. (1979). *How to Share a Secret*. Communications of the ACM. — **Formal (T1)**. Secret sharing scheme used for To key sharding and social recovery.^[?]
27. Heuristic argument. — **Conceptual / speculative (T3)**. The claim that physical-sensor signals are “expensive to simulate” is a cost argument, not a cryptographic proof. Vulnerable to replay, injection, and coercion.^[?]
28. Neulinger, J. & Sparer, K. (2025). *Fostering AI alignment through blockchain, proof of personhood and zero knowledge proofs*. Cluster Computing. <https://link.springer.com/article/10.1007/s10586-025-05729-8> — **Formal / peer-reviewed (T1)**. Establishes PoP consensus, zk-STARK identity verification, AI Shield model. Core engineering reference.^[?]
29. NOSTR Protocol. <https://github.com/nostr-protocol/nips> — **Implementation (T2)**. Decentralized relay protocol. Kind 31990 reserved for application-specific events. Vulnerable to relay censorship and network partition.^[?]
30. Ben-Sasson, E. et al. (2018). *Scalable, transparent, and post-quantum secure computational integrity*. IACR ePrint. — **Formal (T1)**. Foundational STARK paper. Winterfell (Rust) and Cairo are production implementations.^[?]
31. AMD. *SEV-SNP: Strengthening VM Isolation with Integrity Protection and More*. White paper. Intel. *Trust Domain Extensions (TDX) Architecture Specification*. — **Implementation / hardware docs (T2)**. TEEs provide best-available isolation but are vulnerable to side-channels and supply-chain attacks.^[?]
32. Zimmermann, P. (1995). *PGP User’s Guide*. — **Implementation reference (T2)**. Web-of-trust model adapted for biological witness attestation.^[?]
33. TPM / TEE remote attestation literature. — **Implementation analogy (T2)**. Substrate canaries for silicon processes mirror TPM attestation chains but apply to runtime state rather than boot state.^[?]
34. Neulinger, J. & Sparer, K. (2025). *Fostering AI alignment through blockchain, proof of personhood and zero knowledge proofs*. Cluster Computing. <https://link.springer.com/article/10.1007/>

s10586-025-05729-8 — **Formal / peer-reviewed (T1)**. Establishes PoP consensus, zk-STARK identity verification, AI Shield model. Core engineering reference. [?]

35. Neuralink Corporation. *Patient Information: N1 Implant and R1 Robot*. <https://neuralink.com/patient-information/> — **Implementation / medical device (T2)**. As of 2026, Neuralink devices are investigational medical devices with demonstrated thread retraction in early trials. No public remote attestation API exists. All security claims are prospective. [?]
36. Juels, A. & Wattenberg, M. (1999). *A Fuzzy Commitment Scheme*. ACM CCS. — **Formal (T1)**. Allows key derivation from noisy biometrics without storing raw templates. [?]
37. Neuralink Corporation. *Patient Information: N1 Implant and R1 Robot*. <https://neuralink.com/patient-information/> — **Implementation / medical device (T2)**. As of 2026, Neuralink devices are investigational medical devices with demonstrated thread retraction in early trials. No public remote attestation API exists. All security claims are prospective. [?]
38. NeuroLM, SYNAPTICON, Brain-LLM Interface — see [Appendix B](#). **Conceptual / active research (T3)**. Closed-loop EEG→LLM pipelines illustrate sequestration and identity drift threats; not production security claims. [?]
39. Neulinger, J. & Sparer, K. (2025). *Fostering AI alignment through blockchain, proof of personhood and zero knowledge proofs*. Cluster Computing. <https://link.springer.com/article/10.1007/s10586-025-05729-8> — **Formal / peer-reviewed (T1)**. Establishes PoP consensus, zk-STARK identity verification, AI Shield model. Core engineering reference. [?]
40. Ben-Sasson, E. et al. (2018). *Scalable, transparent, and post-quantum secure computational integrity*. IACR ePrint. — **Formal (T1)**. Foundational STARK paper. Winterfell (Rust) and Cairo are production implementations. [?]
41. National Institute of Standards and Technology (NIST). *Post-Quantum Cryptography Standardization*. — **Standards (T1)**. SPHINCS+ selected as hash-based post-quantum signature. Shares hash-function assumptions with zk-STARKs. [?]
42. Ben-Sasson, E. et al. (2018). *Scalable, transparent, and post-quantum secure computational integrity*. IACR ePrint. — **Formal (T1)**. Foundational STARK paper. Winterfell (Rust) and Cairo are production implementations. [?]

43. Juels, A. & Wattenberg, M. (1999). *A Fuzzy Commitment Scheme*. ACM CCS. — **Formal (T1)**. Allows key derivation from noisy biometrics without storing raw templates. [?]
44. Heuristic argument. — **Conceptual / speculative (T3)**. The claim that physical-sensor signals are “expensive to simulate” is a cost argument, not a cryptographic proof. Vulnerable to replay, injection, and coercion. [?]
45. Lamport, L., Shostak, R., & Pease, M. (1982). *The Byzantine Generals Problem*. ACM TOPLAS. — **Formal (T1)**. Establishes the 2/3 consensus bound. No recursive mechanism can exceed this threshold without additional trust assumptions. [?]
46. Neuralink Corporation. *Patient Information: N1 Implant and R1 Robot*. <https://neuralink.com/patient-information/> — **Implementation / medical device (T2)**. As of 2026, Neuralink devices are investigational medical devices with demonstrated thread retraction in early trials. No public remote attestation API exists. All security claims are prospective. [?]
47. Neuralink Corporation. *Patient Information: N1 Implant and R1 Robot*. <https://neuralink.com/patient-information/> — **Implementation / medical device (T2)**. As of 2026, Neuralink devices are investigational medical devices with demonstrated thread retraction in early trials. No public remote attestation API exists. All security claims are prospective. [?]
48. NOSTR Protocol. <https://github.com/nostr-protocol/nips> — **Implementation (T2)**. Decentralized relay protocol. Kind 31990 reserved for application-specific events. Vulnerable to relay censorship and network partition. [?]
49. Neulinger, J. & Sparer, K. (2025). *Fostering AI alignment through blockchain, proof of personhood and zero knowledge proofs*. Cluster Computing. <https://link.springer.com/article/10.1007/s10586-025-05729-8> — **Formal / peer-reviewed (T1)**. Establishes PoP consensus, zk-STARK identity verification, AI Shield model. Core engineering reference. [?]
50. Juels, A. & Wattenberg, M. (1999). *A Fuzzy Commitment Scheme*. ACM CCS. — **Formal (T1)**. Allows key derivation from noisy biometrics without storing raw templates. [?]
51. Ben-Sasson, E. et al. (2018). *Scalable, transparent, and post-quantum secure computational integrity*. IACR ePrint. — **Formal (T1)**. Foundational STARK paper. Winterfell (Rust) and Cairo are production implementations. [?]

52. National Institute of Standards and Technology (NIST). *Post-Quantum Cryptography Standardization*. — **Standards (T1)**. SPHINCS+ selected as hash-based post-quantum signature. Shares hash-function assumptions with zk-STARKs. [?]
53. Shamir, A. (1979). *How to Share a Secret*. Communications of the ACM. — **Formal (T1)**. Secret sharing scheme used for To key sharding and social recovery. [?]
54. AMD. *SEV-SNP: Strengthening VM Isolation with Integrity Protection and More*. White paper. Intel. *Trust Domain Extensions (TDX) Architecture Specification*. — **Implementation / hardware docs (T2)**. TEEs provide best-available isolation but are vulnerable to side-channels and supply-chain attacks. [?]
55. Ben-Sasson, E. et al. (2018). *Scalable, transparent, and post-quantum secure computational integrity*. IACR ePrint. — **Formal (T1)**. Foundational STARK paper. Winterfell (Rust) and Cairo are production implementations. [?]
56. NOSTR Protocol. <https://github.com/nostr-protocol/nips> — **Implementation (T2)**. Decentralized relay protocol. Kind 31990 reserved for application-specific events. Vulnerable to relay censorship and network partition. [?]
57. Neuralink Corporation. *Patient Information: N1 Implant and R1 Robot*. <https://neuralink.com/patient-information/> — **Implementation / medical device (T2)**. As of 2026, Neuralink devices are investigational medical devices with demonstrated thread retraction in early trials. No public remote attestation API exists. All security claims are prospective. [?]
58. Khan, E. (2026). *Brain Hacking: Artificial Intelligence for Safeguarding Against Dangerous AI*. MDPI Preprints. <https://www.preprints.org/manuscript/202601.0156> — **Conceptual / preprint (T3)**. Proposes “cognitive firewall” as neural anomaly detection. Not peer-reviewed. Framing reference only. [?]
59. Banks, I. M. (2010). *Surface Detail*. Orbit. — **Fiction / conceptual inverse (T3)**. Depicts virtual Hells without integrity verification or exit. Used as architectural counter-example, not evidence. [?]
60. Heuristic argument. — **Conceptual / speculative (T3)**. The claim that physical-sensor signals are “expensive to simulate” is a cost argument, not a cryptographic proof. Vulnerable to replay, injection, and coercion. [?]